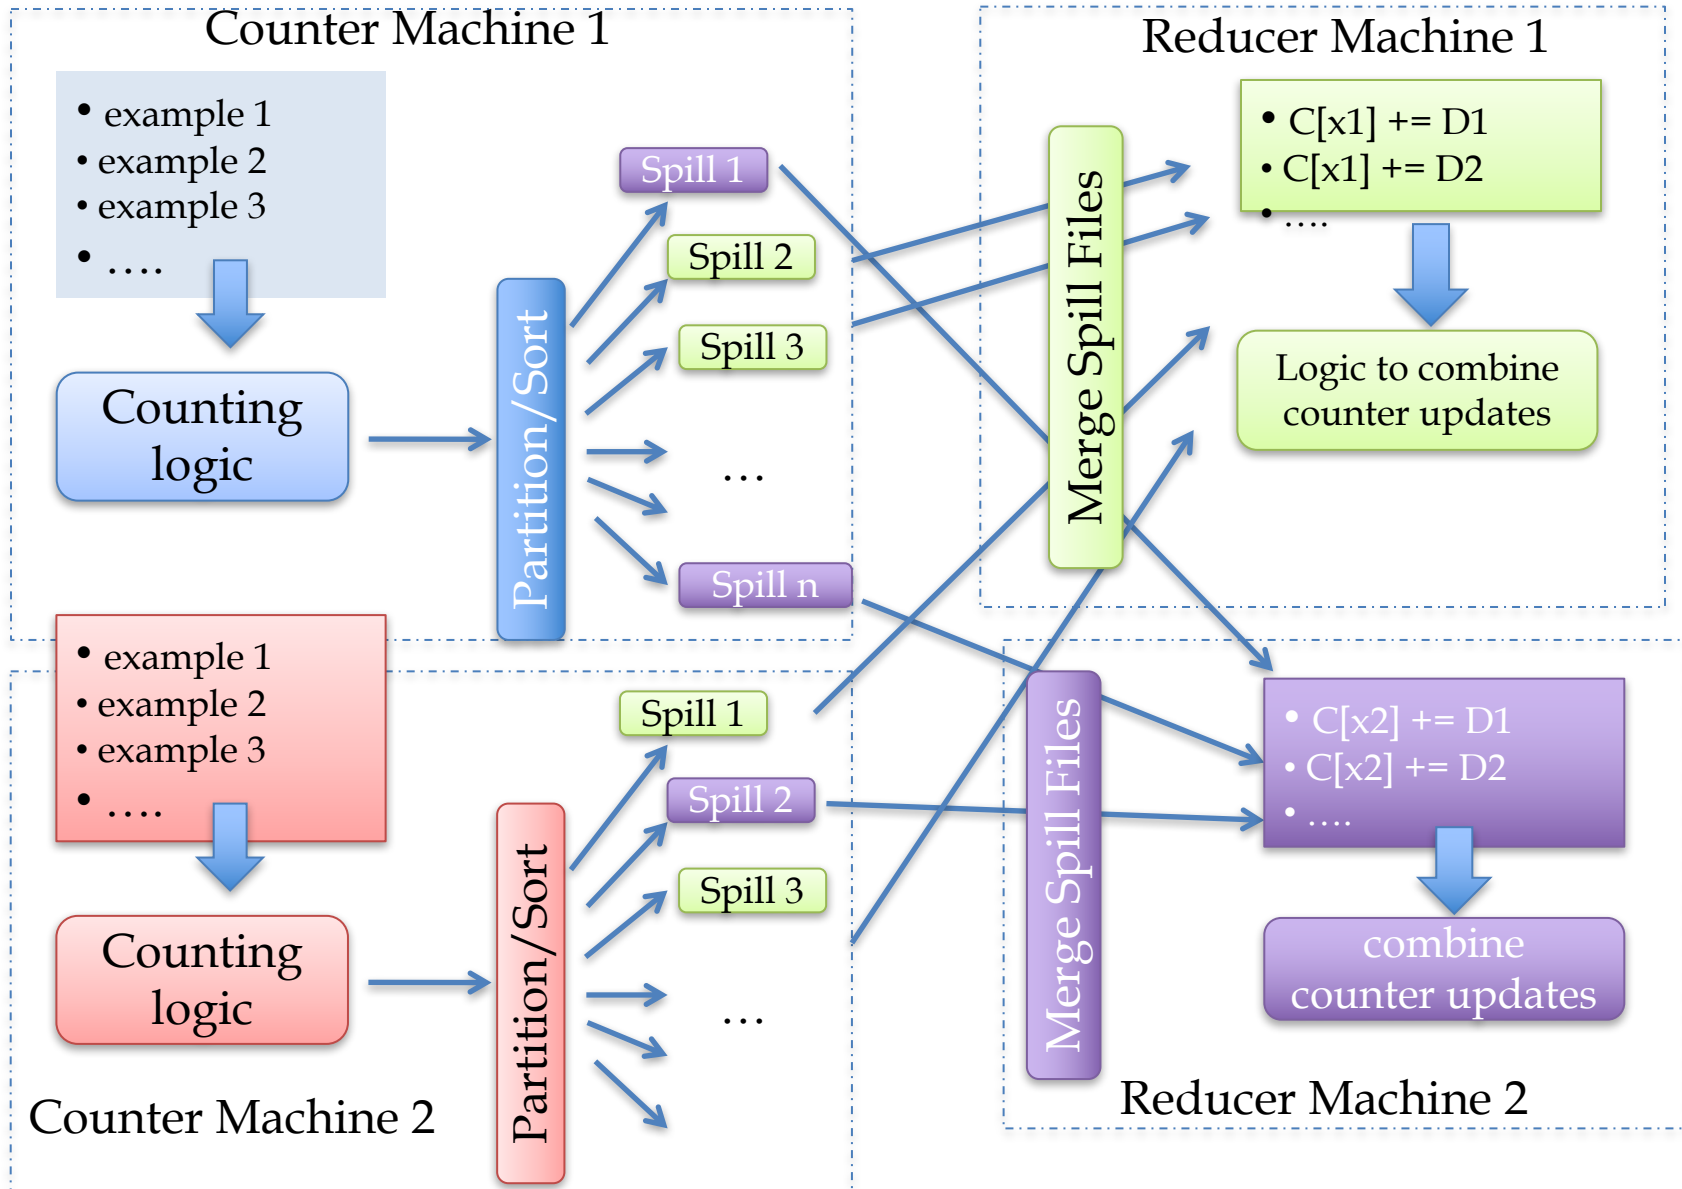


# 10-605: Map-Reduce Workflows

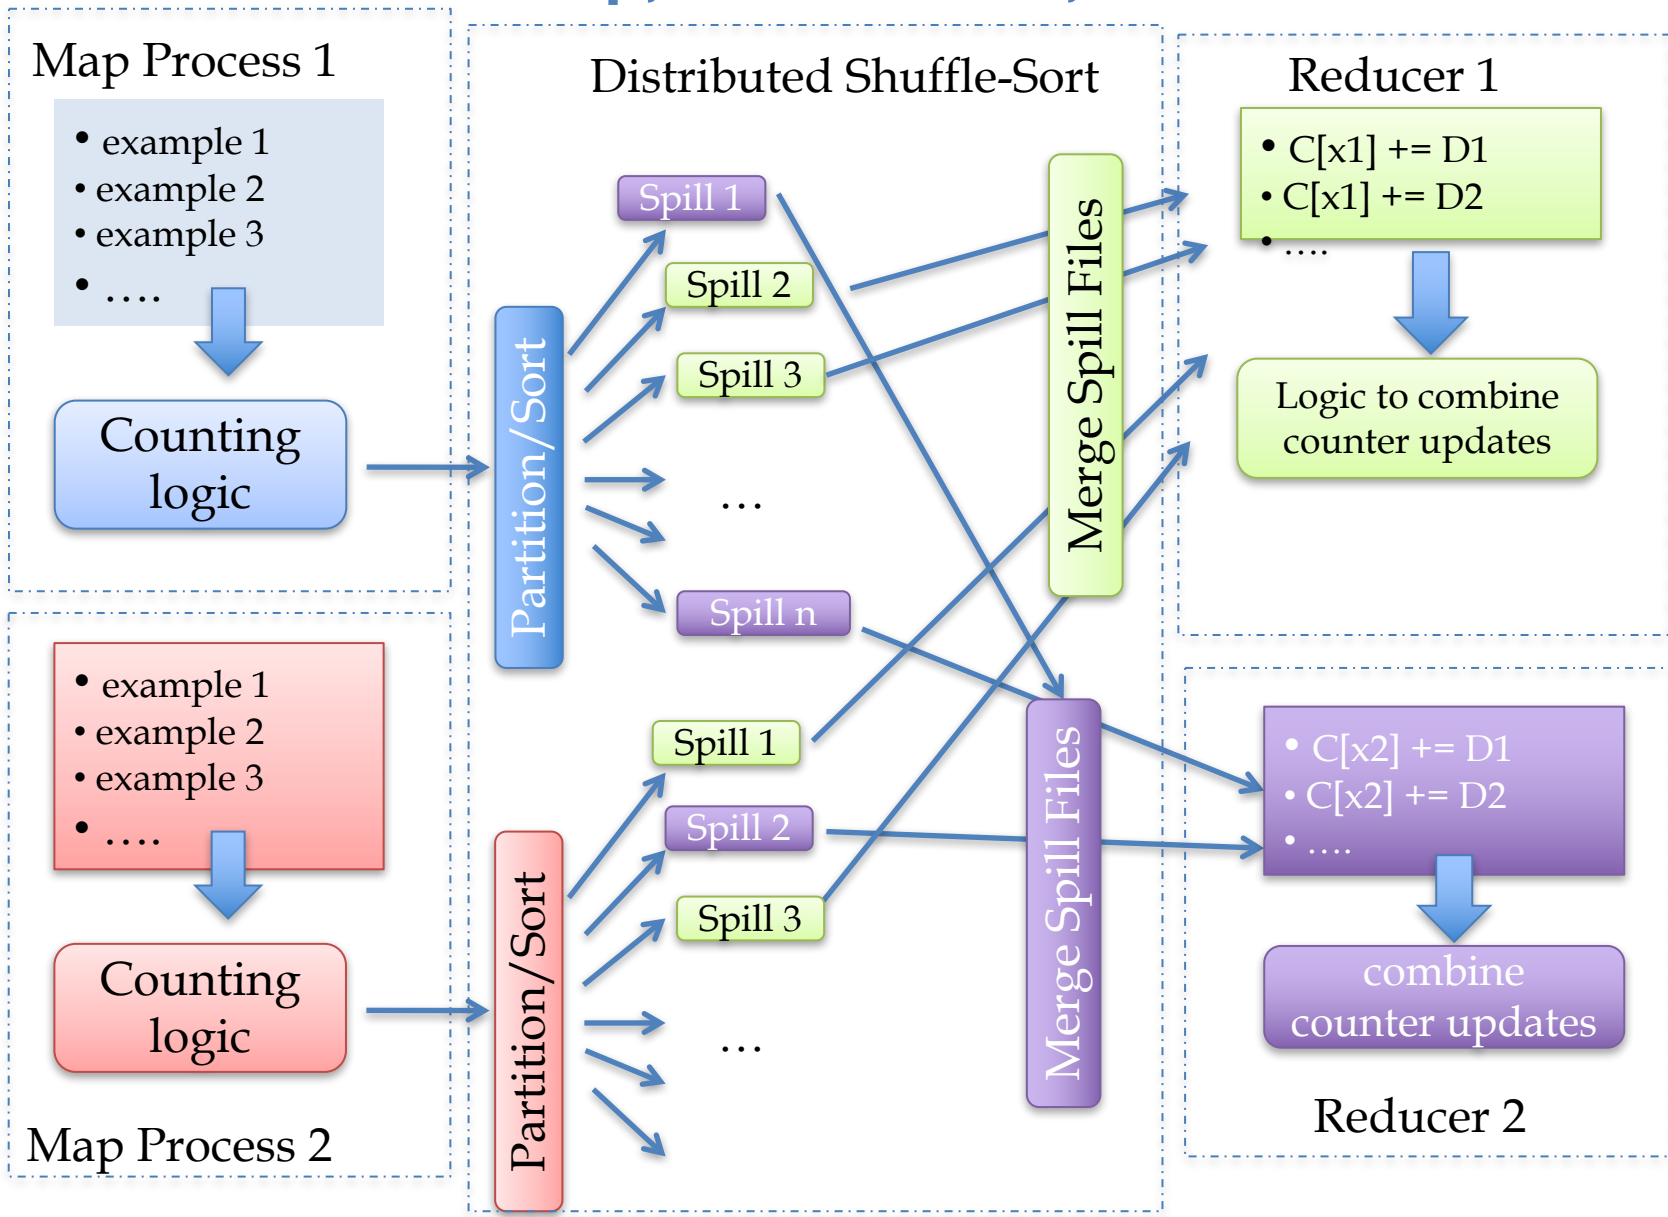
William Cohen

# PARALLELIZING STREAM AND SORT

# Stream and Sort Counting → Distributed Counting



# Distributed Stream-and-Sort: Map, Shuffle-Sort, Reduce



# Combiners in Hadoop

# Some of this is wasteful

- Remember - moving data around and writing to / reading from disk are very expensive operations
- No reducer can start until:
  - all mappers are done
  - data for its partition has been received
  - data in its partition has been sorted

# How much does buffering help?

BUFFER_SIZE	Time	Message Size
<i>none</i>		1.7M <i>words</i>
100	47s	1.2M
1,000	42s	1.0M
10,000	30s	0.7M
100,000	16s	0.24M
1,000,000	13s	0.16M
<i>limit</i>		0.05M

Recall idea here: in stream-and-sort, use a buffer to *accumulate* counts in messages for common words *before* the sort so *sort input was smaller*

# Combiners

- Sits between the map and the shuffle
  - Do some of the reducing while you're waiting for other stuff to happen
  - Avoid moving all of that data over the network
  - Eg, for wordcount: instead of sending (word,1) send (word,n) where n is a *partial count* (over data seen by that mapper)
    - Reducer still just sums the counts
- Only applicable when
  - order of reduce operations doesn't matter (since order is undetermined)
  - effect is cumulative



```

    public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {

private final static IntWritable one = new IntWritable(1);
private Text word = new Text();

    public void map(LongWritable key, Text value, Context context) throws <stuff> {
        String line = value.toString();
        StringTokenizer tokenizer = new StringTokenizer(line);
        while (tokenizer.hasMoreTokens()) {
            word.set(tokenizer.nextToken());
            context.write(word, one);
        }
    }
}

    public static class Reduce extends Reducer<Text, IntWritable, Text,
IntWritable> {

    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}

    public static void main(String[] args) throws Exception {

        Configuration conf = new Configuration();

        Job job = new Job(conf, "wordcount");

        job.setMapperClass(Map.class);
        job.setReducerClass(Reduce.class);

        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);

        job.setOutputKeyClass(Text.class);

        job.setOutputValueClass(IntWritable.class);

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        job.waitForCompletion(true);
    }
}

```

```

public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(LongWritable key, Text value, Context context) throws <stuff> {
        String line = value.toString();
        StringTokenizer tokenizer = new StringTokenizer(line);
        while (tokenizer.hasMoreTokens()) {
            word.set(tokenizer.nextToken());
            context.wr
        }
    }

    public static c
    IntWritable> {

    public void reduc
    throws IOExcept
        int sum = 0;
        for (IntWrita
            sum += va
        }
        context.write
    }
}

    public static void main(String[] args) throws Exception {

        Configuration conf = new Configuration();

        Job job = new Job(conf, "wordcount");

        job.setMapperClass(Map.class);
        job.setReducerClass(Reduce.class);
        job.setCombinerClass(Reduce.class);
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);

        job.setOutputKeyClass(Text.class);

        job.setOutputValueClass(IntWritable.class);

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        job.waitForCompletion(true);
    }
}

```

# Deja vu: Combiner = Reducer

- Often the combiner is the reducer.
  - like for word count
  - but not always
- remember you have no control over when / whether the combiner is applied

# Algorithms in Map-Reduce

(Workflows)

# Scalable out-of-core classification (of large test sets)

can we do better  
that the current approach?

# Testing Large-vocab Naïve Bayes

[For assignment]

- For each example  $id, y, x_1, \dots, x_d$  in *train*:
  - Sort the event-counter update “messages”
  - Scan and add the sorted messages and output the final counter values
- collection of event counts is **big**
- 
- Initialize a HashSet NEEDED and a hashtable C
  - For each example  $id, y, x_1, \dots, x_d$  in *test*:
    - Add  $x_1, \dots, x_d$  to NEEDED
  - For each *event*,  $C(event)$  in the summed counters
    - If *event* involves a NEEDED term  $x$  read it into C
- 
- For each example  $id, y, x_1, \dots, x_d$  in *test*:
    - For each  $y'$  in  $dom(Y)$ :
      - Compute  $\log \Pr(y', x_1, \dots, x_d) = \dots$
- test* is **small**

# Can we do better?

Test data

$id_1$	$w_{1,1}$	$w_{1,2}$	$w_{1,3}$	....	$w_{1,k1}$
$id_2$	$w_{2,1}$	$w_{2,2}$	$w_{2,3}$	....	
$id_3$	$w_{3,1}$	$w_{3,2}$	....		
$id_4$	$w_{4,1}$	$w_{4,2}$	...		
$id_5$	$w_{5,1}$	$w_{5,2}$	....		
..					

Event counts

$X=w_1 \wedge Y=sports$	5245
$X=w_1 \wedge Y=worldNews$	1054
$X=..$	2120
$X=w_2 \wedge Y=...$	37
$X=...$	3
...	...

What we'd like

$id_1$ $w_{1,1}$ $w_{1,2}$ $w_{1,3}$ .... $w_{1,k1}$	$C[X=w_{1,1} \wedge Y=sports]=5245, C[X=w_{1,1} \wedge Y=..], C[X=w_{1,2} \wedge ...]$
$id_2$ $w_{2,1}$ $w_{2,2}$ $w_{2,3}$ ....	$C[X=w_{2,1} \wedge Y=....]=1054, ..., C[X=w_{2,k2} \wedge ...]$
$id_3$ $w_{3,1}$ $w_{3,2}$ ....	$C[X=w_{3,1} \wedge Y=....]=...$
$id_4$ $w_{4,1}$ $w_{4,2}$ ...	...

# Can we do better?

**Step 1:** group counters by word  $w$

How:

- Stream and sort:
  - for each  $C[X=w \wedge Y=y]=n$ 
    - print “ $w \ C[Y=y]=n$ ”
  - sort and build a *list* of values associated with each key  $w$

*Like an inverted index*

Event counts

$X=w_1 \wedge Y=sports$	5245
$X=w_1 \wedge Y=worldNews$	1054
$X=..$	2120
$X=w_2 \wedge Y=...$	37
$X=...$	3
...	...



w	Counts associated with W
aardvark	$C[w \wedge Y=sports]=2$
agent	$C[w \wedge Y=sports]=1027, C[w \wedge Y=worldNews]=564$
...	...
zynga	$C[w \wedge Y=sports]=21, C[w \wedge Y=worldNews]=4464$



# If these records were in a key-value DB we would know what to do....

Test data

$id_1$	$w_{1,1}$	$w_{1,2}$	$w_{1,3}$	....	$w_{1,k1}$
$id_2$	$w_{2,1}$	$w_{2,2}$	$w_{2,3}$	....	
$id_3$	$w_{3,1}$	$w_{3,2}$	....		
$id_4$	$w_{4,1}$	$w_{4,2}$	...		
$id_5$	$w_{5,1}$	$w_{5,2}$	....		
..					

Record of all event counts for each word

w	Counts associated with W
aardvark	$C[w^Y=sports]=2$
agent	$C[w^Y=sports]=1027, C[w^Y=world]$
...	...
zynga	$C[w^Y=sports]=21, C[w^Y=world]$



Step 2: stream through and for each test case

$id_i \ w_{i,1} \ w_{i,2} \ w_{i,3} \ \dots \ w_{i,ki}$

request the event counters needed to classify  $id_i$  from the event-count DB, then classify using the answers



# Is there a stream-and-sort analog of this request-and-answer pattern?

Test data

$id_1$	$w_{1,1}$	$w_{1,2}$	$w_{1,3}$	$\dots$	$w_{1,k1}$
$id_2$	$w_{2,1}$	$w_{2,2}$	$w_{2,3}$	$\dots$	
$id_3$	$w_{3,1}$	$w_{3,2}$	$\dots$		
$id_4$	$w_{4,1}$	$w_{4,2}$	$\dots$		
$id_5$	$w_{5,1}$	$w_{5,2}$	$\dots$		
$\dots$					

Record of all event counts for each word

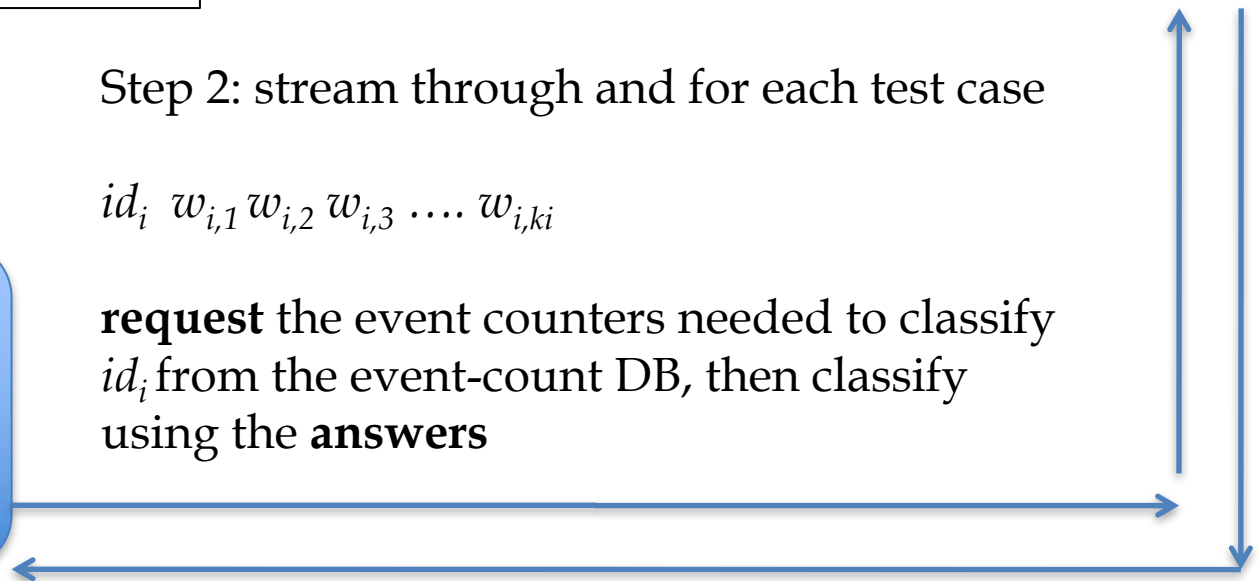
w	Counts associated with W
aardvark	$C[w^Y=sports]=2$
agent	$C[w^Y=sports]=1027, C[w^Y=world]$
...	...
zynga	$C[w^Y=sports]=21, C[w^Y=world]$



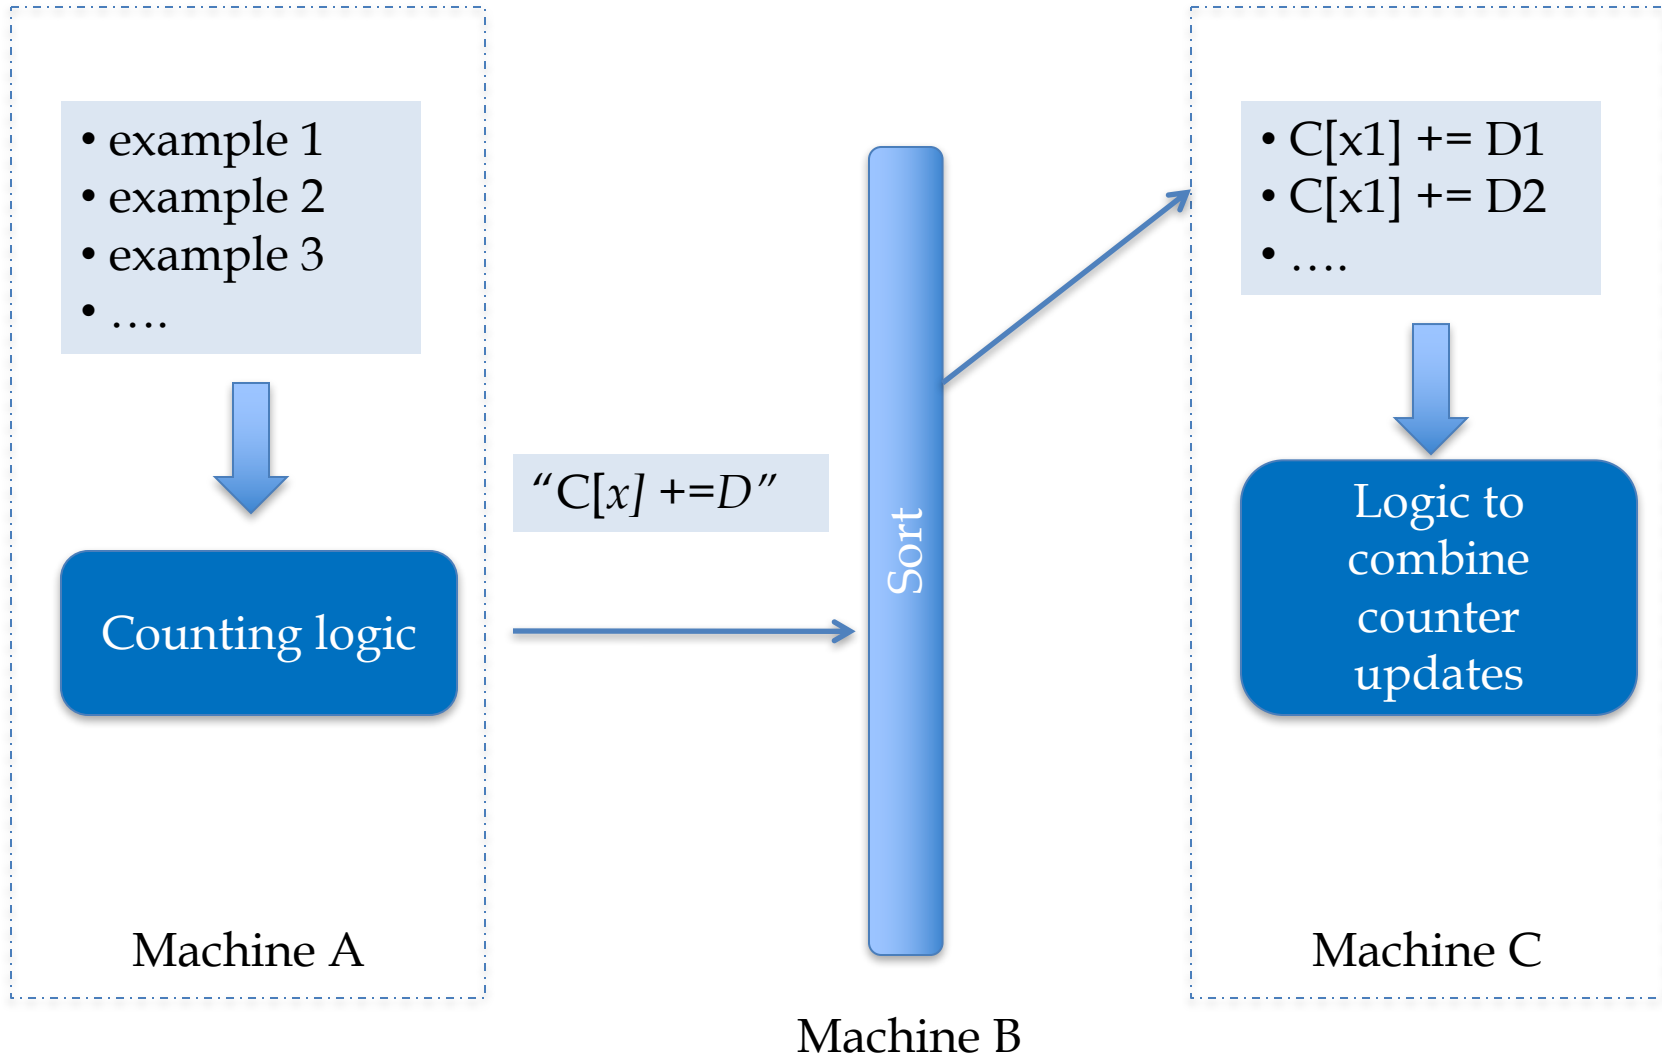
Step 2: stream through and for each test case

$id_i \ w_{i,1} \ w_{i,2} \ w_{i,3} \ \dots \ w_{i,ki}$

**request** the event counters needed to classify  $id_i$  from the event-count DB, then classify using the **answers**



## Recall: Stream and Sort Counting: sort messages so the recipient can stream through them



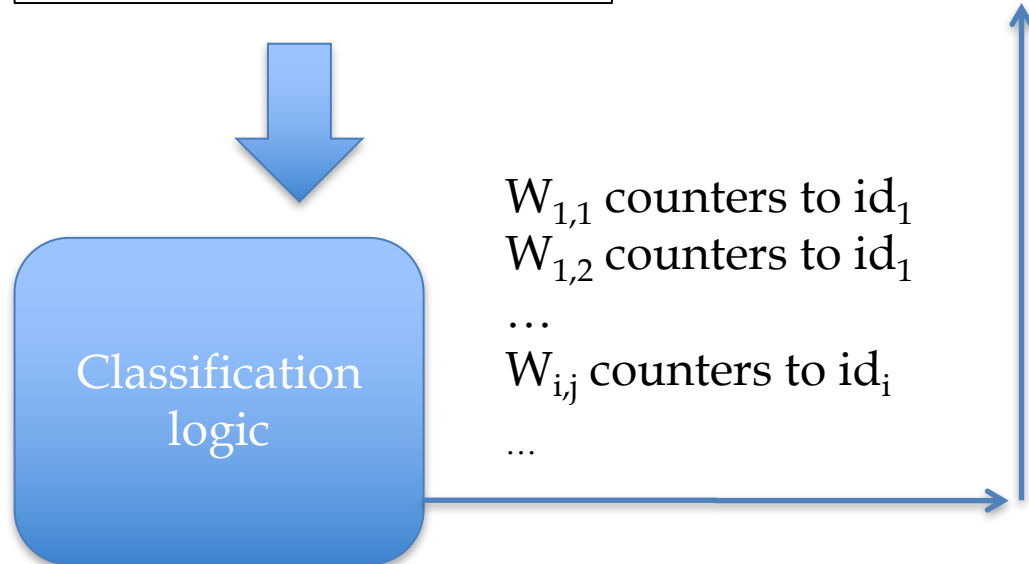
# Is there a stream-and-sort analog of this request-and-answer pattern?

Test data

$id_1$	$w_{1,1}$	$w_{1,2}$	$w_{1,3}$	....	$w_{1,k1}$
$id_2$	$w_{2,1}$	$w_{2,2}$	$w_{2,3}$	....	
$id_3$	$w_{3,1}$	$w_{3,2}$	....		
$id_4$	$w_{4,1}$	$w_{4,2}$	...		
$id_5$	$w_{5,1}$	$w_{5,2}$	....		
..					

Record of all event counts for each word

w	Counts associated with W
aardvark	$C[w^Y=sports]=2$
agent	$C[w^Y=sports]=1027, C[w^Y=world]$
...	...
zynga	$C[w^Y=sports]=21, C[w^Y=world]$



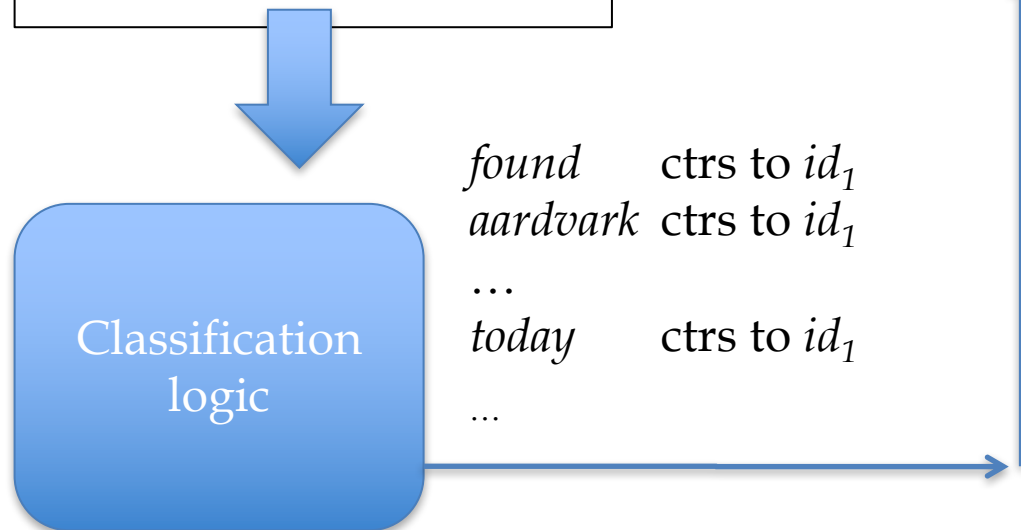
# Is there a stream-and-sort analog of this request-and-answer pattern?

Test data

*id<sub>1</sub> found an aardvark in  
zynga's farmville today!  
id<sub>2</sub> ...  
id<sub>3</sub> ....  
id<sub>4</sub> ...  
id<sub>5</sub> ...  
..*

Record of all event counts for each word

w	Counts associated with W
aardvark	$C[w^Y=\text{sports}]=2$
agent	$C[w^Y=\text{sports}]=1027, C[w^Y=\text{world}]=...$
...	...
zynga	$C[w^Y=\text{sports}]=21, C[w^Y=\text{world}]=...$



# Is there a stream-and-sort analog of this request-and-answer pattern?

Test data

*id<sub>1</sub> found an aardvark in  
zynga's farmville today!  
id<sub>2</sub> ...  
id<sub>3</sub> ....  
id<sub>4</sub> ...  
id<sub>5</sub> ...  
..*

Record of all event counts for each word

w	Counts associated with W
aardvark	C[w^Y=sports]=2
agent	C[w^Y=sports]=1027,C[w^Y=worldN
...	...
zynga	C[w^Y=sports]=21,C[w^Y=worldN



Classification  
logic

*found* ~ctrs to *id<sub>1</sub>*  
*aardvark* ~ctrs to *id<sub>1</sub>*  
...  
*today* ~ctrs to *id<sub>1</sub>*  
...

~ is the last ascii character  
% export LC\_COLLATE=C  
means that it will sort *after*  
anything else with unix sort

# Is there a stream-and-sort analog of this request-and-answer pattern?

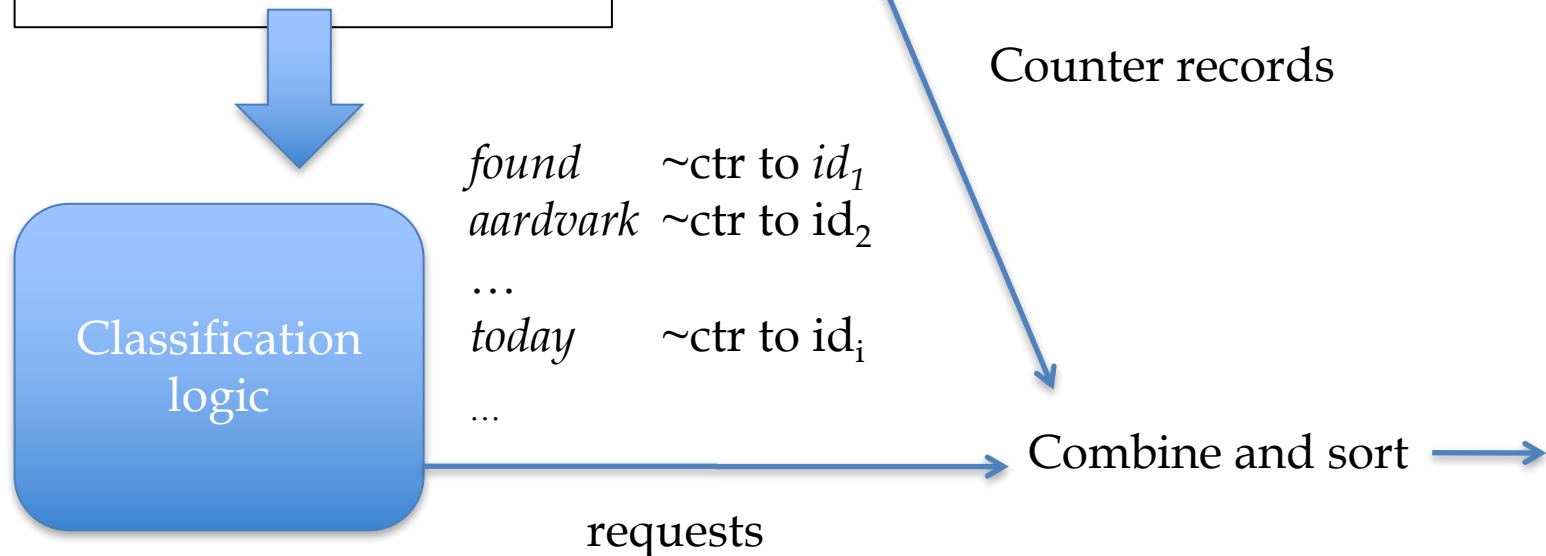
Test data

*id<sub>1</sub> found an aardvark in  
zynga's farmville today!  
id<sub>2</sub> ...  
id<sub>3</sub> ....  
id<sub>4</sub> ...  
id<sub>5</sub> ...  
..*

Record of all event counts for each word

w	Counts associated with W
aardvark	$C[w^Y=\text{sports}]=2$
agent	$C[w^Y=\text{sports}]=1027, C[w^Y=\text{world}]=...$
...	...
zynga	$C[w^Y=\text{sports}]=21, C[w^Y=\text{world}]=...$

Counter records



# A stream-and-sort analog of the request-and-answer pattern...

Record of all event counts for each word

w	Counts
aardvark	$C[w^Y=\text{sports}]=2$
agent	...
...	
zynga	...

Counter records

*found*    ~ctr to  $id_1$   
*aardvark* ~ctr to  $id_1$   
...  
*today*    ~ctr to  $id_1$   
...

Combine and sort

requests

w	Counts
aardvark	$C[w^Y=\text{sports}]=2$
aardvark	~ctr to $id_1$
agent	$C[w^Y=\text{sports}]=...$
agent	~ctr to $id_{345}$
agent	~ctr to $id_{9854}$
...	~ctr to $id_{345}$
agent	~ctr to $id_{34742}$
...	
zynga	$C[...]$
zynga	~ctr to $id_1$

Request-handling logic



# A stream-and-sort analog of the request-and-answer pattern...

- `previousKey = somethingImpossible`
- For each  $(key, val)$  in input:
  - If  $key == previousKey$ 
    - `Answer(recordForPrevKey, val)`
  - Else
    - `previousKey = key`
    - `recordForPrevKey = val`

define `Answer(record, request):`

- find  $id$  where " $request = \sim ctr$  to  $id$ "
- print " $id \sim ctr$  for  $request$  is  $record$ "

w	Counts
aardvark	$C[w^Y = sports] = 2$
aardvark	$\sim ctr$ to id1
agent	$C[w^Y = sports] = \dots$
agent	$\sim ctr$ to id345
agent	$\sim ctr$ to id9854
...	$\sim ctr$ to id345
agent	$\sim ctr$ to id34742
...	
zynga	$C[\dots]$
zynga	$\sim ctr$ to id1

requests → Combine and sort

Request-handling logic

# A stream-and-sort analog of the request-and-answer pattern...

- `previousKey = somethingImpossible`
- For each  $(key, val)$  in input:
  - If  $key == previousKey$ 
    - `Answer(recordForPrevKey, val)`
  - Else
    - `previousKey = key`
    - `recordForPrevKey = val`

define `Answer(record, request):`

- find  $id$  where " $request = \sim ctr \text{ to } id$ "
- print " $id \sim ctr \text{ for } request \text{ is } record$ "

Output:

$id1 \sim ctr \text{ for } aardvark \text{ is } C[w^Y=sports]=2$

...

$id1 \sim ctr \text{ for } zynga \text{ is } ....$

...

Combine and sort

requests

w	Counts
aardvark	$C[w^Y=sports]=2$
aardvark	$\sim ctr \text{ to } id1$
agent	$C[w^Y=sports]=...$
agent	$\sim ctr \text{ to } id345$
agent	$\sim ctr \text{ to } id9854$
...	$\sim ctr \text{ to } id345$
agent	$\sim ctr \text{ to } id34742$
...	
zynga	$C[...]$
zynga	$\sim ctr \text{ to } id1$

Request-handling  
logic

# A stream-and-sort analog of the request-and-answer pattern...

w	Counts
aardvark	$C[w^Y = \text{sports}] = 2$
aardvark	$\sim \text{ctr to id1}$
agent	$C[w^Y = \text{sports}] = \dots$
agent	$\sim \text{ctr to id345}$
agent	$\sim \text{ctr to id9854}$
...	$\sim \text{ctr to id345}$
agent	$\sim \text{ctr to id34742}$
...	
zynga	$C[\dots]$
zynga	$\sim \text{ctr to id1}$

Output:

$id_1 \sim \text{ctr for aardvark is } C[w^Y = \text{sports}] = 2$

...

$id_1 \sim \text{ctr for zynga is } \dots$

...

$id_1$  found an aardvark in  
zynga's farmville today!

$id_2 \dots$

$id_3 \dots$

$id_4 \dots$

$id_5 \dots$

..

Request-handling  
logic

Combine and sort

???? 27

## What we'd wanted

$id_1 \ w_{1,1} w_{1,2} w_{1,3} \dots w_{1,k1}$	$C[X=w_{1,1}^{\wedge}Y=sports]=5245, C[X=w_{1,1}^{\wedge}Y=..], C[X=w_{1,2}^{\wedge}...]$
$id_2 \ w_{2,1} w_{2,2} w_{2,3} \dots$	$C[X=w_{2,1}^{\wedge}Y=....]=1054, ..., C[X=w_{2,k2}^{\wedge}...]$
$id_3 \ w_{3,1} w_{3,2} \dots$	$C[X=w_{3,1}^{\wedge}Y=....]=...$
$id_4 \ w_{4,1} w_{4,2} \dots$	...

## What we ended up with ... and it's good enough!

Key	Value
$id1$	<i>found aardvark zynga farmville today</i>
	$\sim$ ctr for <i>aardvark</i> is $C[w^{\wedge}Y=sports]=2$
	$\sim$ ctr for <i>found</i> is $C[w^{\wedge}Y=sports]=1027, C[w^{\wedge}Y=worldNews]=564$
	...
$id2$	$w_{2,1} w_{2,2} w_{2,3} \dots$
	$\sim$ ctr for $w_{2,1}$ is ...
...	...

# Implementation summary

```
java CountForNB train.dat ... > eventCounts.dat
```

```
java CountsByWord eventCounts.dat | sort
```

```
| java CollectRecords > words.dat
```

```
java requestWordCounts test.dat
```

```
| cat - words.dat | sort | java answerWordCountRequests
```

```
| cat - test.dat | sort | testNBUsingRequests
```

train.dat

$id_1$	$w_{1,1}$	$w_{1,2}$	$w_{1,3}$	....	$w_{1,k1}$
$id_2$	$w_{2,1}$	$w_{2,2}$	$w_{2,3}$	....	
$id_3$	$w_{3,1}$	$w_{3,2}$	....		
$id_4$	$w_{4,1}$	$w_{4,2}$	...		
$id_5$	$w_{5,1}$	$w_{5,2}$	....		
..					

counts.dat

$X=w1^Y=sports$	5245
$X=w1^Y=worldNews$	1054
$X=..$	2120
$X=w2^Y=...$	37
$X=...$	3
...	...

# Implementation summary

```
java CountForNB train.dat ... > eventCounts.dat
```

```
java CountsByWord eventCounts.dat | sort
```

```
| java CollectRecords > words.dat
```

```
java requestWordCounts test.dat
```

```
| cat - words.dat | sort | java answerWordCountRequests
```

```
| cat - test.dat | sort | testNBUsingRequests
```

words.dat

w	Counts associated with W
aardvark	$C[w^Y=\text{sports}]=2$
agent	$C[w^Y=\text{sports}]=1027, C[w^Y=\text{worldNews}]=564$
...	...
zynga	$C[w^Y=\text{sports}]=21, C[w^Y=\text{worldNews}]=4464$

# Implementation summary

```
java CountForNB train.dat ... > eventCounts.dat
java CountsByWord eventCounts.dat | sort
| java CollectRecords > words.dat
```

```
java requestWordCounts test.dat
| cat - words.dat | sort | java answerWordCountRequests
| cat - test.dat | sort | testNBUsingRequests
```

output looks like this

input looks like this

words.dat

*found*      ~ctr to  $id_1$   
*aardvark*   ~ctr to  $id_2$   
 ...  
*today*      ~ctr to  $id_i$   
 ...

w	Counts
aardvark	$C[w^Y=sports]=2$
agent	...
...	
zynga	...

w	Counts
aardvark	$C[w^Y=sports]=2$
aardvark	~ctr to $id_1$
agent	$C[w^Y=sports]=...$
agent	~ctr to $id_{345}$
agent	~ctr to $id_{9854}$
...	~ctr to $id_{345}$

# Implementation summary

```
java CountForNB train.dat ... > eventCounts.dat
java CountsByWord eventCounts.dat | sort
| java CollectRecords > words.dat
```

```
java requestWordCounts test.dat
| cat - words.dat | sort | java answerWordCountRequests
| cat - test.dat | sort | testNBUsingRequests
```

Output  
looks like  
this

Output:

*id1 ~ctr for aardvark is  $C[w^Y = \text{sports}] = 2$*

...

*id1 ~ctr for zynga is ....*

...

**test.dat**

*id<sub>1</sub> found an aardvark in  
zynga's farmville today!*

*id<sub>2</sub> ...*

*id<sub>3</sub> ....*

*id<sub>4</sub> ...*

*id<sub>5</sub> ...*

..



# Implementation summary

```
java CountForNB train.dat ... > eventCounts.dat  
java CountsByWord eventCounts.dat | sort  
| java CollectRecords > words.dat
```

```
java requestWordCounts test.dat  
| cat - words.dat | sort | java answerWordCountRequests  
| cat -test.dat | sort | testNBUsingRequests
```

**Input looks like this**

Key	Value
id1	found aardvark zynga farmville today
	$\sim$ ctr for aardvark is $C[w^Y=\text{sports}]=2$
	$\sim$ ctr for found is $C[w^Y=\text{sports}]=1027, C[w^Y=\text{worldNews}]=564$
	...
id2	$w_{2,1} w_{2,2} w_{2,3} \dots$
	$\sim$ ctr for $w_{2,1}$ is ...
...	...

# ABSTRACTIONS FOR MAP-REDUCE

# Abstractions On Top Of Map-Reduce

- Some obvious streaming processes:
  - for each row in a table
    - Transform it and output the result
  - Decide if you want to keep it with some boolean test, and copy out only the ones that pass the test

**Example:** stem words in a stream of word-count pairs:

$(\text{"aardvarks"}, 1) \rightarrow (\text{"aardvark"}, 1)$

**Proposed syntax:**  $f(\text{row}) \rightarrow \text{row}'$

$\text{table2} = \text{MAP } \text{table1} \text{ TO } \lambda \text{ row} : f(\text{row})$

**Example:** apply stop words

$(\text{"aardvark"}, 1) \rightarrow (\text{"aardvark"}, 1)$   
 $(\text{"the"}, 1) \rightarrow \text{deleted}$

**Proposed syntax:**  $f(\text{row}) \rightarrow \{\text{true}, \text{false}\}$

$\text{table2} = \text{FILTER } \text{table1} \text{ BY } \lambda \text{ row} : f(\text{row})$

# Abstractions On Top Of Map-Reduce

- A non-obvious? streaming processes:
  - for each row in a table
    - Transform it to a list of items
    - Splice all the lists together to get the output table (**flatten**)

Proposed syntax:

$f(row) \rightarrow \text{list of rows}$

`table2 = FLATMAP table1 TO  $\lambda row: f(row)$`

**Example:** tokenizing a line

“I found an aardvark”  $\rightarrow$  [“i”, “found”, “an”, “aardvark”]

“We love zymurgy”  $\rightarrow$  [“we”, “love”, “zymurgy”]

..but final table is one word per row

“i”

“found”

“an”

“aardvark”

“we”

“love”

...

# NB Test Step

How:

- Stream and sort:
    - for each  $C[X=w \wedge Y=y]=n$ 
      - print “w C[Y=y]=n”
    - sort and build a *list* of values associated with each key  $w$
- Like an inverted index*

Event counts

$X=w_1 \wedge Y=sports$	5245
$X=w_1 \wedge Y=worldNews$	1054
$X=..$	2120
$X=w_2 \wedge Y=...$	37
$X=...$	3
...	...



w	Counts associated with W
aardvark	$C[w \wedge Y=sports]=2$
agent	$C[w \wedge Y=sports]=1027, C[w \wedge Y=worldNews]=564$
...	...
zynga	$C[w \wedge Y=sports]=21, C[w \wedge Y=worldNews]=4464$

# NB Test Step

## The general case:

We're taking rows from a table

- In a particular format (*event,count*)

Applying a function to get a new value

- The *word* for the event

And *grouping* the rows of the table by this new value

## → Grouping operation

*Special case of a map-reduce*

## Event counts

$X=w_1 \wedge Y=sports$	5245
$X=w_1 \wedge Y=worldNews$	1054
$X=..$	2120
$X=w_2 \wedge Y=...$	37
$X=...$	3
...	...



**Proposed syntax:**  $f(row) \rightarrow field$

GROUP *table* BY  $\lambda row: f(row)$

Could define  $f$  via: a function, a field of a defined *record* structure, ...

w	Counts associated with W
aardvark	$C[w \wedge Y=sports]=2$
agent	$C[w \wedge Y=sports]=1027, C[w \wedge Y=worldNews]=564$
...	...
zynga	$C[w \wedge Y=sports]=21, C[w \wedge Y=worldNews]=4464$

# NB Test Step

## The general case:

We're taking rows from a table

- In a particular format (*event,count*)

Applying a function to get a new value

- The *word* for the event

And *grouping* the rows of the table by this new value

## → Grouping operation

*Special case of a map-reduce*

Aside: you guys know how to implement this, right?

1. Output pairs  $(f(\text{row}), \text{row})$  with a map/streaming process
2. Sort pairs by key – which is  $f(\text{row})$
3. Reduce and aggregate by *appending together* all the values associated with the same key

**Proposed syntax:**  $f(\text{row}) \rightarrow \text{field}$

GROUP *table* BY  $\lambda \text{ row} : f(\text{row})$

Could define  $f$  via: a function, a field of a defined *record* structure, ...

# Abstractions On Top Of Map-Reduce

- And another example from the Naïve Bayes test program...



# Request-and-answer

Test data

$id_1$	$w_{1,1}$	$w_{1,2}$	$w_{1,3}$	....	$w_{1,k1}$
$id_2$	$w_{2,1}$	$w_{2,2}$	$w_{2,3}$	....	
$id_3$	$w_{3,1}$	$w_{3,2}$	....		
$id_4$	$w_{4,1}$	$w_{4,2}$	...		
$id_5$	$w_{5,1}$	$w_{5,2}$	....		
..					

Record of all event counts for each word

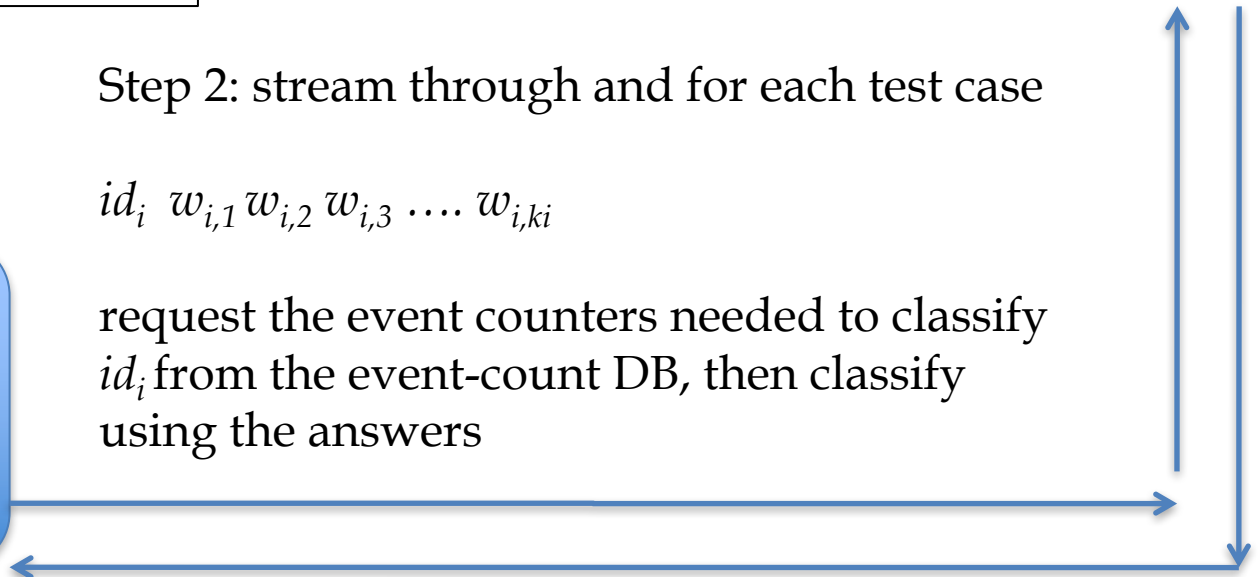
w	Counts associated with W
aardvark	$C[w^Y=sports]=2$
agent	$C[w^Y=sports]=1027, C[w^Y=world]$
...	...
zynga	$C[w^Y=sports]=21, C[w^Y=world]$



Step 2: stream through and for each test case

$id_i \ w_{i,1} \ w_{i,2} \ w_{i,3} \ \dots \ w_{i,ki}$

request the event counters needed to classify  $id_i$  from the event-count DB, then classify using the answers



# Request-and-answer

- Break down into stages
  - Generate the data being requested (indexed by key, here a word)
    - Eg with group ... by
  - Generate the requests as (key, requestor) pairs
    - Eg with flatmap ... to
  - **Join** these two tables by key
    - Join: conceptually defined as (1) cross-product and (2) filter out pairs with different values for keys
    - Join: implemented by concatenating two different tables of key-value pairs, and reducing them together
  - Postprocess the joined result

w	Request
found	~ctr to id1
aardvark	~ctr to id1
...	
zynga	~ctr to id1
...	~ctr to id2

w	Counters
aardvark	$C[w^Y=\text{sports}]=2$
agent	$C[w^Y=\text{sports}]=1027, C[w^Y=\text{worldNews}]$
...	...
zynga	$C[w^Y=\text{sports}]=21, C[w^Y=\text{worldNews}]$

w	Counters	Requests
aardvark	$C[w^Y=\text{sports}]=2$	~ctr to id1
agent	$C[w^Y=\text{sports}]=...$	~ctr to id345
agent	$C[w^Y=\text{sports}]=...$	~ctr to id9854
agent	$C[w^Y=\text{sports}]=...$	~ctr to id345
...	$C[w^Y=\text{sports}]=...$	~ctr to id34742
zynga	$C[...]$	~ctr to id1
zynga	$C[...]$	...

w	Request
found	id1
aardvark	id1
...	
zynga	id1
...	id2

w	Counters
aardvark	$C[w^Y=sports]=2$
agent	$C[w^Y=sports]=1027, C[w^Y=worldNews]=...$
...	...
zynga	$C[w^Y=sports]=21, C[w^Y=worldNews]=...$

## Examples:

JOIN *wordInDoc* BY *word*, *wordCounters* BY *word* --- if *word(row)* defined correctly

JOIN *wordInDoc* BY lambda (word,docid):word,  
*wordCounters* BY lambda (word,counters):word – using python syntax for functions

## Proposed syntax:

JOIN *table1* BY  $\lambda row: f(row)$ ,  
*table2* BY  $\lambda row: g(row)$

$C[w^Y=sports]=...$	id345
$C[w^Y=sports]=...$	id34742
$C[...]$	id1
$C[...]$	...

# Abstract Implementation: [TF]IDF

1/2

data = pairs (docid, term) where term is a word appears in document with id docid

opera

- DIS
- GRO

key	value
found	(d123,found),(d134,found),... 2456
aardvark	(d123,aardvark),... 7

ce st

docId	term
d123	found
d123	aardvark

docFreq = DISTINCT data

| GROUP BY  $\lambda(\text{docid}, \text{term}): \text{term}$  REDUCING TO

key	value
1	12451

docIds = MAP DATA BY  $\lambda(\text{docid}, \text{term}): \text{docid}$  | DISTINCT

numDocs = GROUP docIds BY  $\lambda \text{docid}: 1$  REDUCING TO count /\* (1, numDocs) \*/

dataPlusDF =

question – how many reducers should I use here?

JOIN data BY  $\lambda(\text{docid}, \text{term}): \text{term}$ , docFreq BY  $\lambda(\text{term}, \text{df}): \text{term}$

| MAP  $\lambda((\text{docid}, \text{term}), (\text{term}, \text{df})):(\text{docId}, \text{term}, \text{df})$  /\* (docId, term, document-freq) \*/

unnormalizedDocVecs = JOIN dataPlusDF by  $\lambda \text{row}: 1$ , numDocs by  $\lambda \text{row}: 1$

| MAP  $\lambda((\text{docId}, \text{term}, \text{df}), (\text{dummy}, \text{numDocs})):(\text{docId}, \text{term}, \log(\text{numDocs} / \text{df}))$

/\* (docId, term, weight-before-normalizing) : u \*/

# Abstract Implementation: [TF]IDF

1/2

data = pairs (docid, term) where term is a word appears in document with id docid

operation

key	value
found	(d123,found),(d134,found),... 2456
aardvark	(d123,aardvark),... 7

• DIS

• GROUP

ce st

docId	term
d123	found
d123	aardvark

docFreq = DISTINCT data

| GROUP BY  $\lambda(\text{docid}, \text{term}): \text{term}$  REDUCING TO

key	value
1	12451

docIds = MAP DATA BY  $\lambda(\text{docid}, \text{term}): \text{docid}$  | DISTINCT

numDocs = GROUP docIds BY  $\lambda \text{docid}: 1$  REDUCING TO count /\* (1, numDocs) \*/

dataPlusDF =

question – how many reducers should I use here?

JOIN data BY  $\lambda(\text{docid}, \text{term}): \text{term}$ , docFreq BY  $\lambda(\text{term}, \text{df}): \text{term}$

| MAP  $\lambda((\text{docid}, \text{term}), (\text{term}, \text{df})): (\text{docId}, \text{term}, \text{df})$  /\* (docId, term, document-freq) \*/

unnormalizedDocVecs = JOIN dataPlusDF by  $\lambda \text{row}: 1$ , numDocs by  $\lambda \text{row}: 1$

| MAP  $\lambda((\text{docId}, \text{term}, \text{df}), (\text{dummy}, \text{numDocs})): (\text{docId}, \text{term}, \log(\text{numDocs} / \text{df}))$

/\* (docId, term, weight-before-normalizing) · 1 \*/

question – how many reducers should I use here?

# Abstract Implementation: TFIDF

2/2

normalizers =

```
GROUP unnormlizedDocVecs BY  $\lambda(\text{docId}, \text{term}, w): \text{docId}$   
RETAINING  $\lambda(\text{docId}, \text{term}, w): w^2$   
REDUCING TO sum /* (docId, sum-of-square-weights) */
```

key	
d1234	(d1234,found,1.542), (d1234,aardvark,13.23),... 37.234
d3214	.... 29.654

```
docVec = JOIN unnormlizedDocVecs BY  $\lambda(\text{docId}, \text{term}, w): \text{docId}$ ,  
          normalizers BY  $\lambda(\text{docId}, \text{norm}): \text{docId}$   
  | MAP  $\lambda((\text{docId}, \text{term}, w), (\text{docId}, \text{norm})): (\text{docId}, \text{term}, w / \sqrt{\text{norm}})$   
  /* (docId, term, weight) */
```

docId	term	w	docId	w
d1234	found	1.542	d1234	37.234
d1234	aardvark	13.23	d1234	37.234

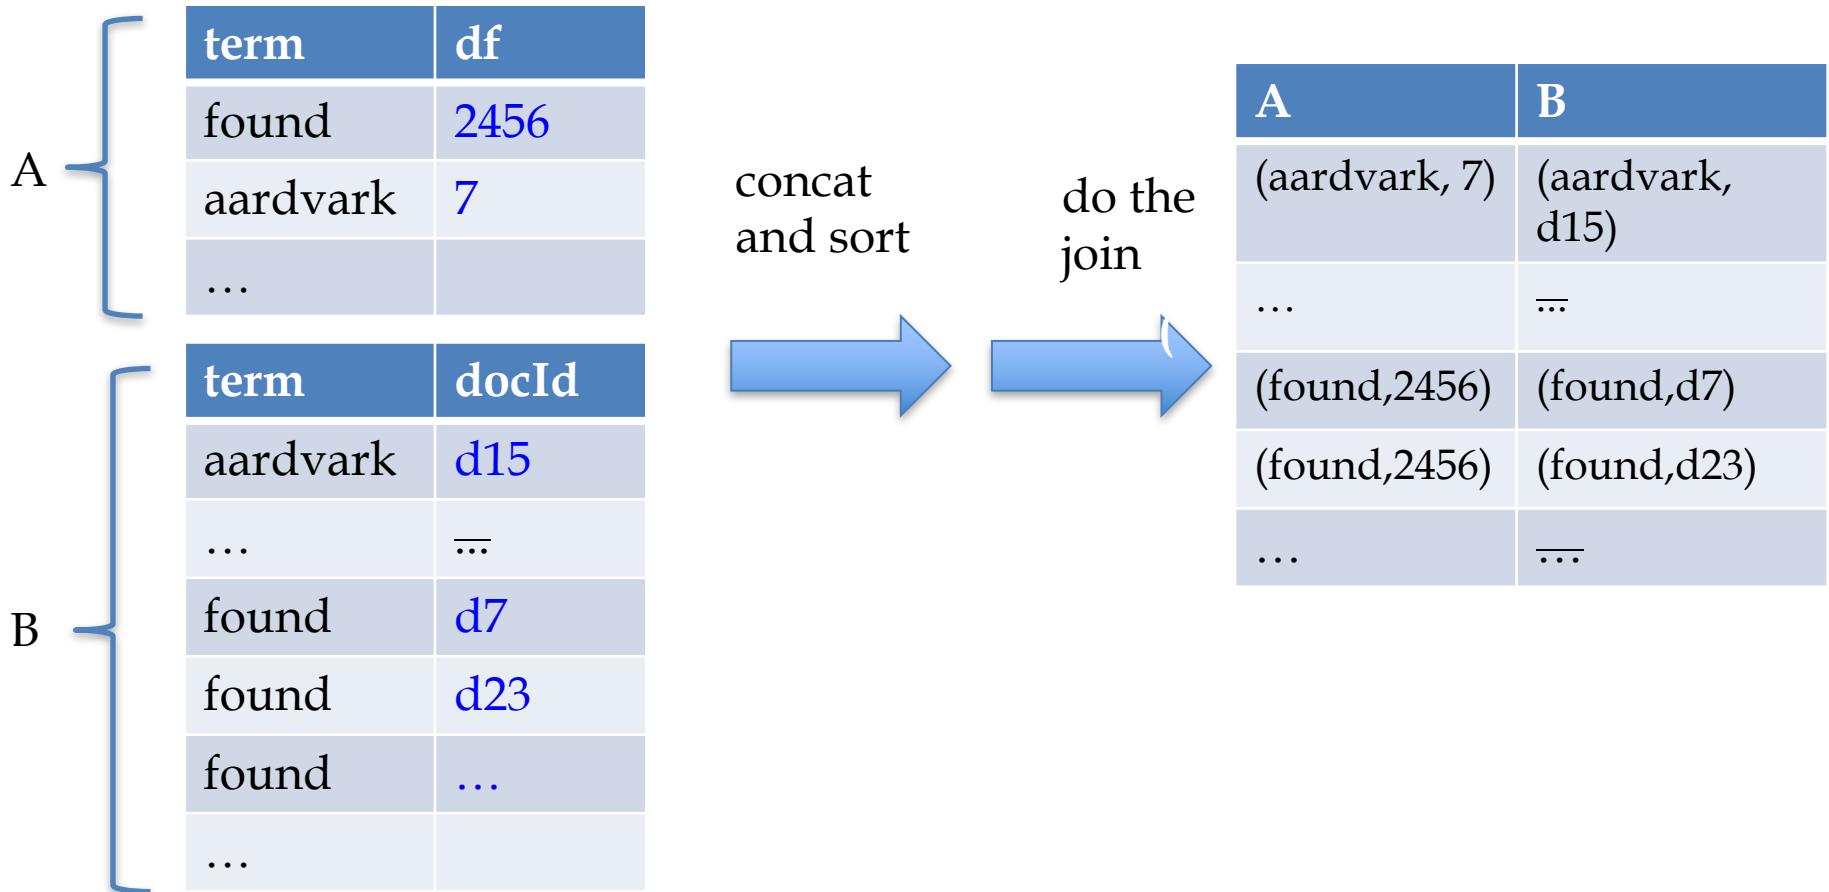
# Two ways to join

- Reduce-side join
- Map-side join



# Two ways to join

- Reduce-side join for A,B




# Two ways to join

- Reduce-side join for A,B

	df
	2456
ark	7

	docId
ark	d15
	...
	d7
	d23
	...

concat  
and sort



term		df
found	A	2456
aardvark	A	7
...		

term		docId
aardvark	B	d15
...		...
found	B	d7
found	B	d23
found	B	...
...		

do the  
join



A	B
(aardvark, 7)	(aardvark, d15)
...	...
(found, 2456)	(found, d7)
(found, 2456)	(found, d23)
...	...

tricky bit: need **sort** by **first two** values (aardvark, AB) – we want the DF's to come first

but all tuples with key “aardvark” should go to **same worker**


# Two ways to join

- Reduce-side join for A,B

	df
	2456
ark	7

concat  
and sort

	docId
ark	d15
	...
	d7
	d23
	...



term		df
found	A	2456
aardvark	A	7
...		

term		docId
aardvark	B	d15
...		...
found	B	d7
found	B	d23
found	B	...
...		

tricky bit: need **sort** by **first two** values (aardvark, AB) – we want the DF's to come first

but all tuples with key "aardvark" should go to **same worker**

**custom sort**  
**(secondary sort key):**  
**Writable with your**  
**own Comparator**

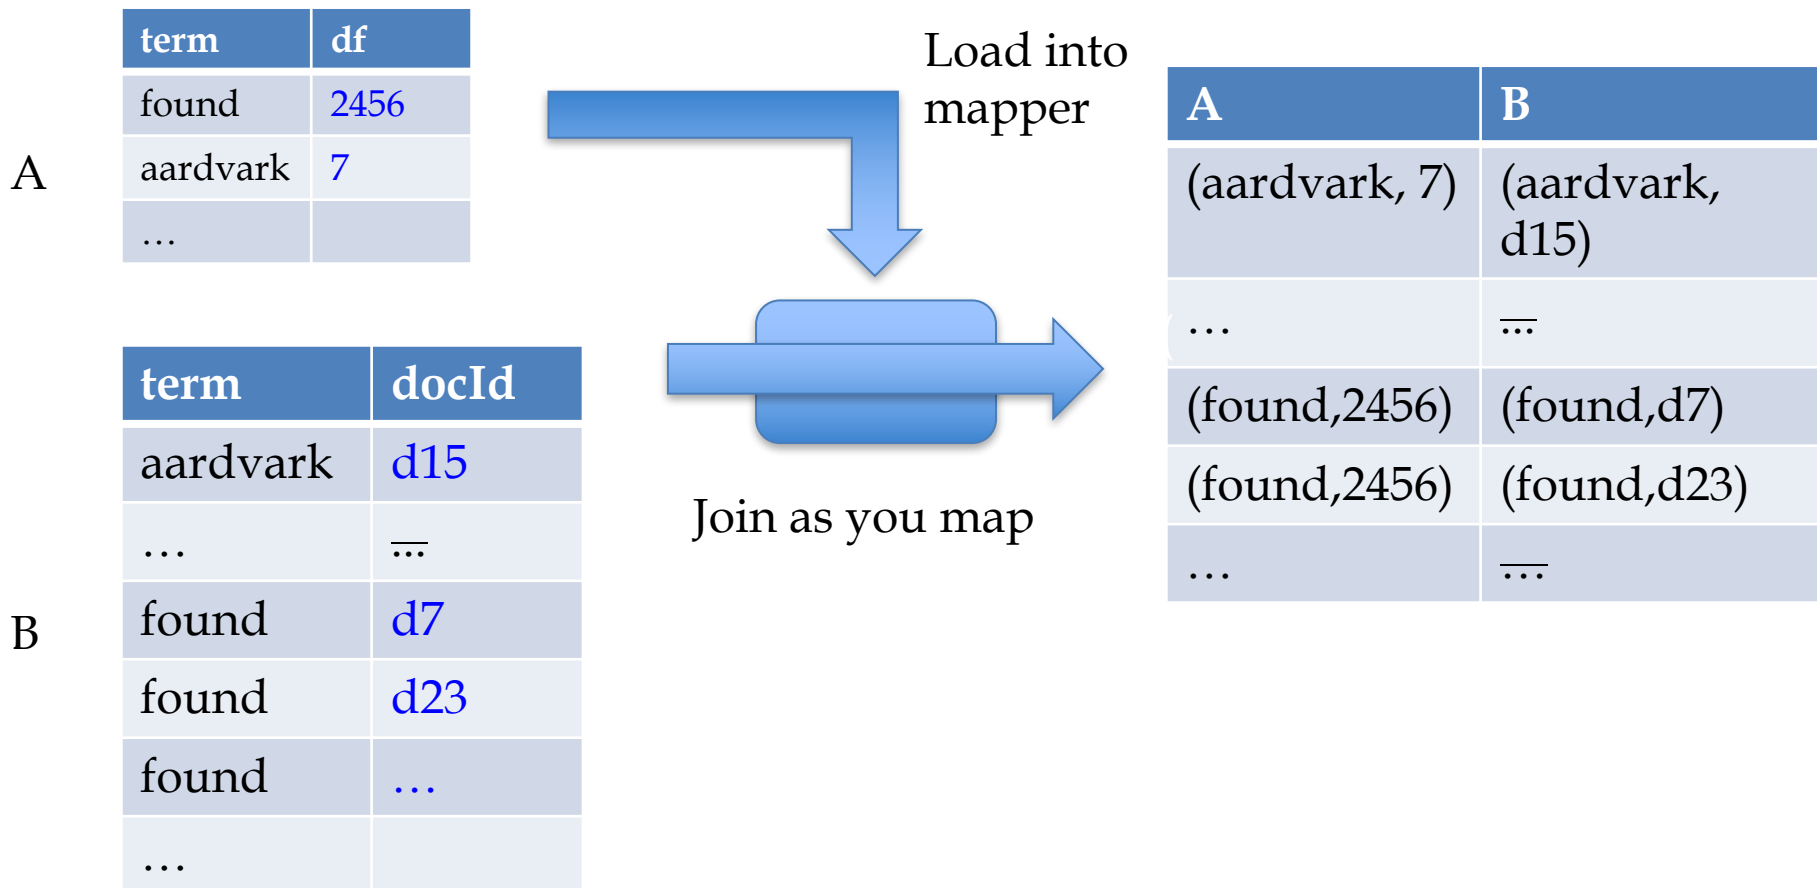
**custom Partitioner**  
**(specified for job**  
**like the Mapper,**  
**Reducer, ..)**

# Two ways to join

- Map-side join
  - write the smaller relation out to disk
  - send it to each Map worker
    - DistributedCache
  - when you **initialize** each Mapper, load in the small relation
    - Configure(...) is called at initialization time
  - map through the larger relation and do the join
  - faster but requires one relation to go in memory

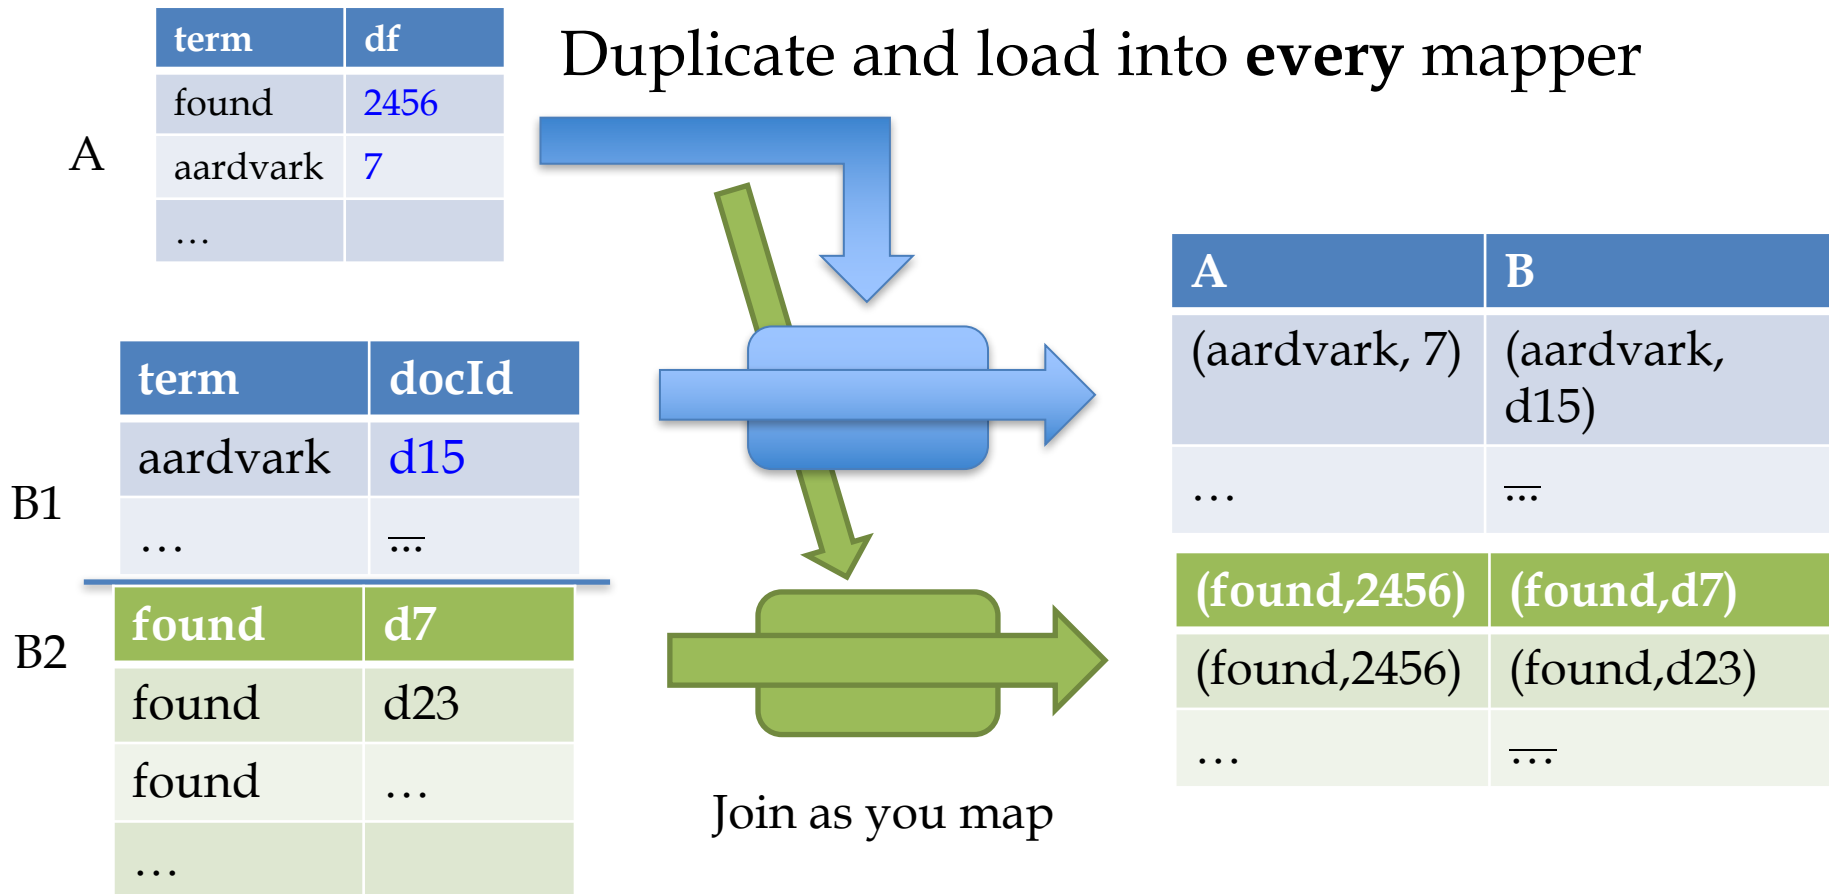
# Two ways to join

- Map-side join for A (small) and B (large)



# Two ways to join

- Map-side join for A (small) and B (large)



# **FIG: A WORKFLOW LANGUAGE**

# PIG: word count example

```
A = load '/tmp/bible+shakes.nopunc';  
B = foreach A generate flatten(TOKENIZE((chararray)$0)) as word;  
C = filter B by word matches '\w+';  
D = group C by word;  
E = foreach D generate COUNT(C) as count, group as word;  
F = order E by count desc;  
store F into '/tmp/wc';
```

PIG program is a bunch of **assignments** where every LHS is a **relation**.

No loops, conditionals, etc allowed.



```
A = load '/tmp/bible+shakes.nopunc';  
B = foreach A generate flatten(TOKENIZE((chararray)$0)) as word;  
C = filter B by word matches '\w+';  
D = group C by word;  
E = foreach D generate COUNT(C) as count, group as word;  
F = order E by count desc;  
store F into '/tmp/wc';
```

Tokenize – built-in function

Built-in regex matching ....

Flatten – special keyword, which applies to the **next** step in the process – so output is a stream of words w/o document boundaries

```
A = load '/tmp/bible+shakes.nopunc';  
B = foreach A generate flatten(TOKENIZE((chararray)$0)) as word;  
C = filter B by word matches '\w+';  
D = group C by word;  
E = foreach D generate COUNT(C) as count, group as word;  
F = order E by count desc;  
store F into '/tmp/wc';
```

Group produces a stream of **bags** of identical words... bags, tuples, dictionaries are primitive types

Group by ... foreach generate count(...) will be optimized into a single map-reduce