

CMU SCS

15-826: Multimedia Databases and Data Mining

Lecture #27: Time series mining and
forecasting
Christos Faloutsos




CMU SCS

Must-Read Material


- Byong-Kee Yi, Nikolaos D. Sidiropoulos, Theodore Johnson, H.V. Jagadish, Christos Faloutsos and Alex Biliris, *Online Data Mining for Co-Evolving Time Sequences*, ICDE, Feb 2000.
- Chungmin Melvin Chen and Nick Roussopoulos, *Adaptive Selectivity Estimation Using Query Feedbacks*, SIGMOD 1994

15-826 (c) C. Faloutsos, 2017 2




CMU SCS


Thanks



Deepay Chakrabarti (UT-Austin)



Spiros Papadimitriou (Rutgers)



Prof. Byoung-Kee Yi (Samsung)

15-826 (c) C. Faloutsos, 2017 3



CMU SCS

Outline

➔

- Motivation
 - Similarity search – distance functions
 - Linear Forecasting
 - Bursty traffic - fractals and multifractals
 - Non-linear forecasting
- Conclusions

15-826 (c) C. Faloutsos, 2017 4

CMU SCS

Problem definition

- Given: one or more sequences
 $x_1, x_2, \dots, x_t, \dots$
 $(y_1, y_2, \dots, y_p, \dots)$
 \dots)
- Find
 - similar sequences; forecasts
 - patterns; clusters; outliers

15-826 (c) C. Faloutsos, 2017 5

CMU SCS

Motivation - Applications

- Financial, sales, economic series
- Medical
 - ECGs +; blood pressure etc monitoring
 - reactions to new drugs
 - elderly care

15-826 (c) C. Faloutsos, 2017 6

CMU SCS

Motivation - Applications (cont' d)

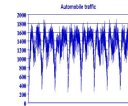
- 'Smart house'
 - sensors monitor temperature, humidity, air quality
- video surveillance

15-826 (c) C. Faloutsos, 2017 7

CMU SCS

Motivation - Applications (cont' d)

- civil/automobile infrastructure
 - bridge vibrations [Oppenheim+02]
 - road conditions / traffic monitoring

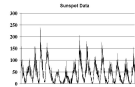


15-826 (c) C. Faloutsos, 2017 8

CMU SCS

Motivation - Applications (cont'd)

- Weather, environment/anti-pollution
 - volcano monitoring
 - air/water pollutant monitoring



15-826 (c) C. Faloutsos, 2017 9

CMU SCS

Motivation - Applications (cont'd)

- Computer systems
 - ‘Active Disks’ (buffering, prefetching)
 - web servers (ditto)
 - network traffic monitoring
 - ...

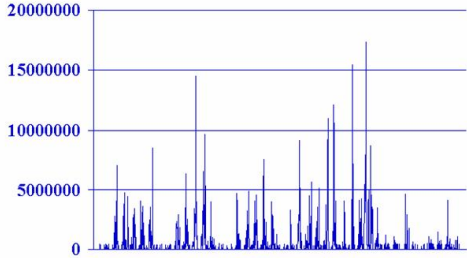
15-826 (c) C. Faloutsos, 2017 10

CMU SCS

Stream Data: Disk accesses

#bytes

Disk traffic



time

15-826 (c) C. Faloutsos, 2017 11

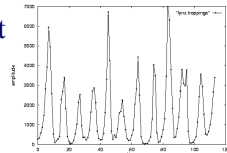
CMU SCS

Problem #1:

Goal: given a signal (e.g., #packets over time)

Find: patterns, periodicities, and/or compress





count



lynx caught per year
(packets per day;
temperature per day)

year

15-826 (c) C. Faloutsos, 2017 12



CMU SCS

Outline

- Motivation
- ➔ • Similarity Search and Indexing
- Linear Forecasting
- Bursty traffic - fractals and multifractals
- Non-linear forecasting
- Conclusions

15-826 (c) C. Faloutsos, 2017 17

CMU SCS

Outline

- Motivation
- ➔ • Similarity search and distance functions
 - Euclidean
 - Time-warping
- ...

15-826 (c) C. Faloutsos, 2017 18

CMU SCS

Importance of distance functions

Subtle, but **absolutely necessary**:

- A ‘must’ for similarity indexing (-> forecasting)
- A ‘must’ for clustering

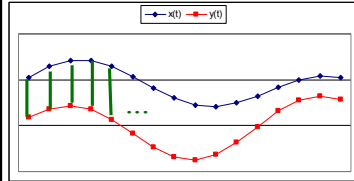
Two major families

- Euclidean and Lp norms
- Time warping and variations

15-826 (c) C. Faloutsos, 2017 19

CMU SCS

Euclidean and Lp



$$D(\vec{x}, \vec{y}) = \sum_{i=1}^n (x_i - y_i)^2$$

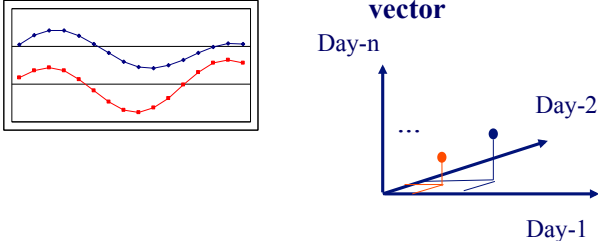
$$L_p(\vec{x}, \vec{y}) = \sum_{i=1}^n |x_i - y_i|^p$$

- L_1 : city-block = Manhattan
- L_2 = Euclidean
- L_∞

15-826 (c) C. Faloutsos, 2017 20

Observation #1

- Time sequence \rightarrow n-d vector

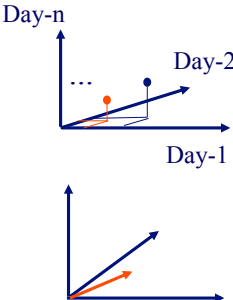


15-826 (c) C. Faloutsos, 2017 21

Observation #2

Euclidean distance is closely related to

- cosine similarity
- dot product
- ‘cross-correlation’ function



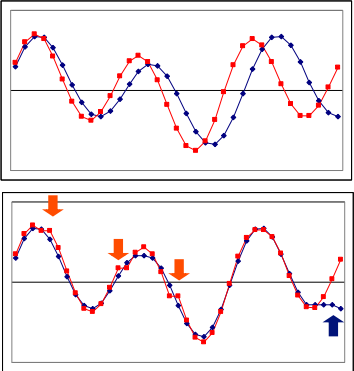
15-826 (c) C. Faloutsos, 2017 22

Time Warping

- allow accelerations - decelerations
 - (with or w/o penalty)
- THEN compute the (Euclidean) distance (+ penalty)
- related to the string-editing distance


15-826 (c) C. Faloutsos, 2017 23

Time Warping



‘stutters’ :


15-826 (c) C. Faloutsos, 2017 24



Time warping

Q: how to compute it?
 A: dynamic programming
 $D(i, j)$ = cost to match
 prefix of length i of first sequence x with prefix
 of length j of second sequence y


15-826 (c) C. Faloutsos, 2017 25



Full-text scanning

- Approximate matching - **string editing** distance:
 $d(\text{'survey'}, \text{'surgery'}) = 2$
 = min # of insertions, deletions, substitutions to transform the first string into the second
 SURVEY
 SURGERY

15-826 (c) C. Faloutsos, 2017 26




Time warping

Thus, with no penalty for stutter, for sequences
 $x_1, x_2, \dots, x_i, \dots, y_1, y_2, \dots, y_j$

$$D(i, j) = \|x[i] - y[j]\| + \min \begin{cases} D(i-1, j-1) & \text{no stutter} \\ D(i, j-1) & \text{x-stutter} \\ D(i-1, j) & \text{y-stutter} \end{cases}$$

15-826 (c) C. Faloutsos, 2017 27



Time warping

VERY SIMILAR to the string-editing distance

$$D(i, j) = \|x[i] - y[j]\| + \min \begin{cases} D(i-1, j-1) & \text{no stutter} \\ D(i, j-1) & \text{x-stutter} \\ D(i-1, j) & \text{y-stutter} \end{cases}$$

15-826 (c) C. Faloutsos, 2017 28

CMU SCS

Full-text scanning

```

if s[i] = t[j] then
  cost(i, j) = cost(i-1, j-1)
else
  cost(i, j) = min (
    1 + cost(i, j-1) // deletion
    1 + cost(i-1, j-1) // substitution
    1 + cost(i-1, j) // insertion
  )

```

15-826 (c) C. Faloutsos, 2017 29

CMU SCS

Time warping

VERY SIMILAR to the string-editing distance

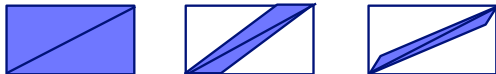
Time-warping	String editing
$D(i, j) = \ x[i] - y[j]\ + \min \begin{cases} D(i-1, j-1) \\ D(i, j-1) \\ D(i-1, j) \end{cases}$	$\text{cost}(i, j) = \min \begin{cases} 1 + \text{cost}(i-1, j-1) // \text{sub.} \\ 1 + \text{cost}(i, j-1) // \text{del.} \\ 1 + \text{cost}(i-1, j) // \text{ins.} \end{cases}$

15-826 (c) C. Faloutsos, 2017 30

CMU SCS

Time warping

- Complexity: $O(M*N)$ - quadratic on the length of the strings
- Many** variations (penalty for stutters; limit on the number/percentage of stutters; ...)
- popular in voice processing [Rabiner + Juang]



15-826 (c) C. Faloutsos, 2017 31


CMU SCS

Other Distance functions

- piece-wise linear/flat approx.; compare pieces [Keogh+01] [Faloutsos+97]
- 'cepstrum' (for voice [Rabiner+Juang])
 - do DFT; take log of amplitude; do DFT again!
- Allow for small gaps [Agrawal+95]

See tutorial by [Gunopulos + Das, SIGMOD01]

15-826 (c) C. Faloutsos, 2017 32




CMU SCS

Other Distance functions

- In [Keogh+, KDD' 04]: parameter-free, MDL based

15-826 (c) C. Faloutsos, 2017 33




CMU SCS

Conclusions

Prevailing distances:

- Euclidean and
- time-warping

15-826 (c) C. Faloutsos, 2017 34




CMU SCS

Outline

- Motivation
- Similarity search and distance functions
- ➔ • Linear Forecasting
- Bursty traffic - fractals and multifractals
- Non-linear forecasting
- Conclusions

15-826 (c) C. Faloutsos, 2017 35



CMU SCS

Linear Forecasting

15-826 (c) C. Faloutsos, 2017 36

CMU SCS

Forecasting

"Prediction is very difficult, especially about the future." - Nils Bohr

<http://www.hfac.uh.edu/MediaFutures/thoughts.html>

15-826 (c) C. Faloutsos, 2017 37

CMU SCS

Outline

- Motivation
- ...
- Linear Forecasting
 - Auto-regression: Least Squares; RLS
 - Co-evolving time sequences
 - Examples
 - Conclusions

15-826 (c) C. Faloutsos, 2017 38

CMU SCS

Reference

[Yi+00] Byoung-Kee Yi et al.: *Online Data Mining for Co-Evolving Time Sequences*, ICDE 2000. (Describes MUSCLES and Recursive Least Squares)

15-826 (c) C. Faloutsos, 2017 39

CMU SCS

Problem#2: Forecast

- Example: give x_{t-1}, x_{t-2}, \dots , forecast x_t

15-826 (c) C. Faloutsos, 2017 40

CMU SCS

Forecasting: Preprocessing

MANUALLY:

- remove trends
- spot periodicities

7 days

15-826 (c) C. Faloutsos, 2017 41

CMU SCS

Problem#2: Forecast

- Solution: try to express x_t as a linear function of the past: x_{t-2}, x_{t-2}, \dots (up to a window of w)

Formally:

$$x_t \approx a_1 x_{t-1} + \dots + a_w x_{t-w} + noise$$

15-826 (c) C. Faloutsos, 2017 42

CMU SCS

(Problem: Back-cast; interpolate)

- Solution - interpolate: try to express x_t as a linear function of the past AND the future: $x_{t+1}, x_{t+2}, \dots, x_{t+w_{future}}; x_{t-1}, \dots, x_{t-w_{past}}$ (up to windows of w_{past}, w_{future})
- EXACTLY the same algo's

15-826 (c) C. Faloutsos, 2017 43

CMU SCS

Linear Regression: idea

patient	weight	height
1	27	43
2	43	54
3	54	72
...
N	25	??

- express what we don't know (= 'dependent variable')
- as a linear function of what we know (= 'indep. variable(s)')

15-826 (c) C. Faloutsos, 2017 44

CMU SCS

Linear Auto Regression:

Time	Packets Sent(t)
1	43
2	54
3	72
...	...
N	??

15-826 (c) C. Faloutsos, 2017 45

CMU SCS

Linear Auto Regression:

Time	Packets Sent ($t-1$)	Packets Sent(t)
1	-	43
2	43	54
3	54	72
...
N	25	??

- lag $w=1$
- Dependent variable = # of packets sent ($S[t]$)
- Independent variable = # of packets sent ($S[t-1]$)

15-826 (c) C. Faloutsos, 2017 46

CMU SCS

Outline

- Motivation
- ...
- Linear Forecasting
 - Auto-regression: **Least Squares; RLS**
 - Co-evolving time sequences
 - Examples
 - Conclusions

15-826 (c) C. Faloutsos, 2017 47

CMU SCS

More details:

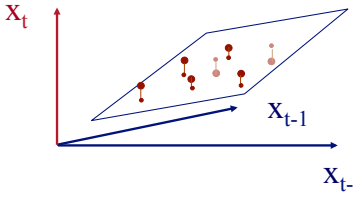
- Q1: Can it work with window $w>1$?
- A1: YES!

15-826 (c) C. Faloutsos, 2017 48

CMU SCS

More details:

- Q1: Can it work with window $w > 1$?
- A1: YES! (we'll fit a hyper-plane, then!)

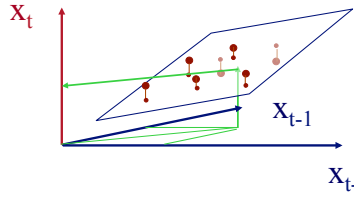


15-826 (c) C. Faloutsos, 2017 49

CMU SCS

More details:

- Q1: Can it work with window $w > 1$?
- A1: YES! (we'll fit a hyper-plane, then!)



15-826 (c) C. Faloutsos, 2017 50

CMU SCS

More details:

- Q1: Can it work with window $w > 1$?
- A1: YES! The problem becomes:

$$\mathbf{X}_{[N \times w]} \times \mathbf{a}_{[w \times 1]} = \mathbf{y}_{[N \times 1]}$$

- OVER-CONSTRAINED
 - \mathbf{a} is the vector of the regression coefficients
 - \mathbf{X} has the N values of the w indep. variables
 - \mathbf{y} has the N values of the dependent variable

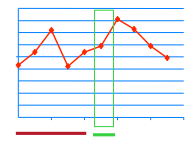
15-826 (c) C. Faloutsos, 2017 51

CMU SCS

More details:

- $\mathbf{X}_{[N \times w]} \times \mathbf{a}_{[w \times 1]} = \mathbf{y}_{[N \times 1]}$

Ind-var1 Ind-var-w



time

$$\begin{bmatrix} X_{11}, X_{12}, \dots, X_{1w} \\ X_{21}, X_{22}, \dots, X_{2w} \\ \vdots \\ X_{N1}, X_{N2}, \dots, X_{Nw} \end{bmatrix} \times \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_w \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}$$

15-826 (c) C. Faloutsos, 2017 52

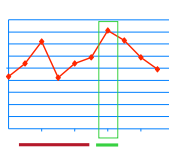
CMU SCS

More details:

- $\mathbf{X}_{[N \times w]} \times \mathbf{a}_{[w \times 1]} = \mathbf{y}_{[N \times 1]}$

Ind-var1 Ind-var-w

time



$$\begin{bmatrix} X_{11}, X_{12}, \dots, X_{1w} \\ X_{21}, X_{22}, \dots, X_{2w} \\ \vdots \\ X_{N1}, X_{N2}, \dots, X_{Nw} \end{bmatrix} \times \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_w \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}$$

15-826 (c) C. Faloutsos, 2017 53

CMU SCS

More details

- Q2: How to estimate $a_1, a_2, \dots, a_w = \mathbf{a}$?
- A2: with Least Squares fit

$$\mathbf{a} = (\mathbf{X}^T \times \mathbf{X})^{-1} \times (\mathbf{X}^T \times \mathbf{y})$$

- (Moore-Penrose pseudo-inverse)
- \mathbf{a} is the vector that minimizes the RMSE from \mathbf{y}
- <identical math with 'query feedbacks'>

15-826 (c) C. Faloutsos, 2017 54

CMU SCS

More details

- Q2: How to estimate $a_1, a_2, \dots, a_w = \mathbf{a}$?
- A2: with Least Squares fit

$$\mathbf{a} = (\mathbf{X}^T \times \mathbf{X})^{-1} \times (\mathbf{X}^T \times \mathbf{y})$$

Identical to earlier formula (proof?)

$$\mathbf{a} = \mathbf{V} \times \mathbf{\Lambda}^{(-1)} \times \mathbf{U}^T \times \mathbf{y}$$

Where

$$\mathbf{X} = \mathbf{U} \times \mathbf{\Lambda} \times \mathbf{V}^T$$

15-826 (c) C. Faloutsos, 2017 55

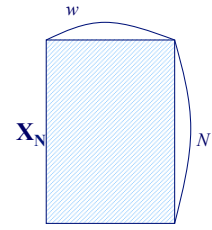
CMU SCS

More details

- Straightforward solution:

$$\mathbf{a} = (\mathbf{X}^T \times \mathbf{X})^{-1} \times (\mathbf{X}^T \times \mathbf{y})$$

\mathbf{a} : Regression Coeff. Vector
 \mathbf{X} : Sample Matrix



- Observations:
 - Sample matrix \mathbf{X} grows over time
 - needs matrix inversion
 - $\mathcal{O}(N \times w^2)$ computation
 - $\mathcal{O}(N \times w)$ storage

15-826 (c) C. Faloutsos, 2017 56

CMU SCS

Even more details

- Q3: Can we estimate \mathbf{a} incrementally?
- A3: Yes, with the brilliant, classic method of 'Recursive Least Squares' (RLS) (see, e.g., [Yi+00], for details).
- We can do the matrix inversion, WITHOUT inversion! (How is that possible?!)

15-826 (c) C. Faloutsos, 2017 57

CMU SCS

Even more details

- Q3: Can we estimate \mathbf{a} incrementally?
- A3: Yes, with the brilliant, classic method of 'Recursive Least Squares' (RLS) (see, e.g., [Yi+00], for details).
- We can do the matrix inversion, WITHOUT inversion! (How is that possible?!)
- A: our matrix has special form: $(\mathbf{X}^T \mathbf{X})$

15-826 (c) C. Faloutsos, 2017 58

CMU SCS

More details

At the $N+1$ time tick:

15-826 (c) C. Faloutsos, 2017 59

CMU SCS

More details

$$\mathbf{a} = (\mathbf{X}^T \times \mathbf{X})^{-1} \times (\mathbf{X}^T \times \mathbf{y})$$


- Let $\mathbf{G}_N = (\mathbf{X}_N^T \times \mathbf{X}_N)^{-1}$ ('gain matrix')
- \mathbf{G}_{N+1} can be computed recursively from \mathbf{G}_N

15-826 (c) C. Faloutsos, 2017 60

CMU SCS

EVEN more details:

$$G_{N+1} = G_N - [c]^{-1} \times [G_N \times x_{N+1}^T] \times x_{N+1} \times G_N$$



 $l \times w$ row vector

$$c = [1 + x_{N+1} \times G_N \times x_{N+1}^T]$$

Let's elaborate
(VERY IMPORTANT, VERY VALUABLE!)

15-826 (c) C. Faloutsos, 2017 61

CMU SCS

EVEN more details:

$$a = [X_{N+1}^T \times X_{N+1}]^{-1} \times [X_{N+1}^T \times y_{N+1}]$$

15-826 (c) C. Faloutsos, 2017 62

CMU SCS

EVEN more details:

$$a = [X_{N+1}^T \times X_{N+1}]^{-1} \times [X_{N+1}^T \times y_{N+1}]$$

$[w \times 1]$ $[(N+1) \times w]$ $[(N+1) \times 1]$
 $[w \times (N+1)]$ $[w \times (N+1)]$

15-826 (c) C. Faloutsos, 2017 63

CMU SCS

EVEN more details:

$$a = [X_{N+1}^T \times X_{N+1}]^{-1} \times [X_{N+1}^T \times y_{N+1}]$$

$[(N+1) \times w]$
 $[w \times (N+1)]$

15-826 (c) C. Faloutsos, 2017 64

CMU SCS

EVEN more details:

$$a = [X_{N+1}^T \times X_{N+1}]^{-1} \times [X_{N+1}^T \times y_{N+1}]$$

‘gain matrix’, $G_{N+1} \equiv [X_{N+1}^T \times X_{N+1}]^{-1}$ 1 x w row vector

$$G_{N+1} = G_N - [c]^{-1} \times [G_N \times x_{N+1}^T] \times x_{N+1} \times G_N$$

$$c = [1 + x_{N+1} \times G_N \times x_{N+1}^T]$$

15-826 (c) C. Faloutsos, 2017 65

CMU SCS

EVEN more details:

$$G_{N+1} = G_N - [c]^{-1} \times [G_N \times x_{N+1}^T] \times x_{N+1} \times G_N$$

$$c = [1 + x_{N+1} \times G_N \times x_{N+1}^T]$$

15-826 (c) C. Faloutsos, 2017 66

CMU SCS

EVEN more details:

1x1

wxw wxw wxw wx1 wxw

$$G_{N+1} = G_N - [c]^{-1} \times [G_N \times x_{N+1}^T] \times x_{N+1} \times G_N$$

SCALAR! $c = [1 + x_{N+1} \times G_N \times x_{N+1}^T]$

15-826 (c) C. Faloutsos, 2017 67

CMU SCS

$a = (X^T \times X)^{-1} \times (X^T \times y)$

Altogether:

$$a = [X_{N+1}^T \times X_{N+1}]^{-1} \times [X_{N+1}^T \times y_{N+1}]$$

$$G_{N+1} \equiv [X_{N+1}^T \times X_{N+1}]^{-1}$$

$$G_{N+1} = G_N - [c]^{-1} \times [G_N \times x_{N+1}^T] \times x_{N+1} \times G_N$$

$$c = [1 + x_{N+1} \times G_N \times x_{N+1}^T]$$

15-826 (c) C. Faloutsos, 2017 68

CMU SCS

Altogether:

$$G_0 \equiv \delta I$$

IMPORTANT!

where
 I : $w \times w$ identity matrix
 δ : a large positive number (say, 10^4)

15-826 (c) C. Faloutsos, 2017 69

CMU SCS

Comparison:

- **Straightforward Least Squares**
 - Needs huge matrix (**growing** in size) $O(N \times w)$
 - Costly matrix operation $O(N \times w^2)$
- **Recursive LS**
 - Need much smaller, fixed size matrix $O(w \times w)$
 - Fast, incremental computation $O(1 \times w^2)$
 - **no matrix inversion**

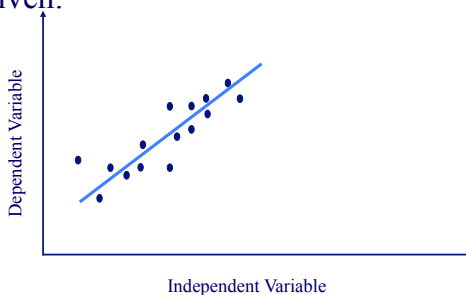
$N = 10^6, \quad w = 1-100$

15-826 (c) C. Faloutsos, 2017 70

CMU SCS

Pictorially:

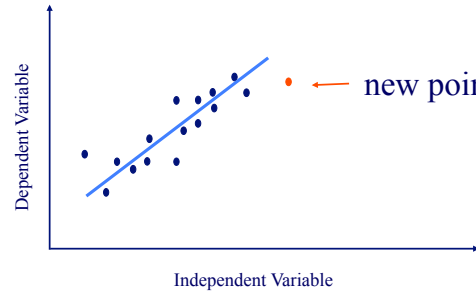
- **Given:**




15-826 (c) C. Faloutsos, 2017 71

CMU SCS

Pictorially:



15-826 (c) C. Faloutsos, 2017 72

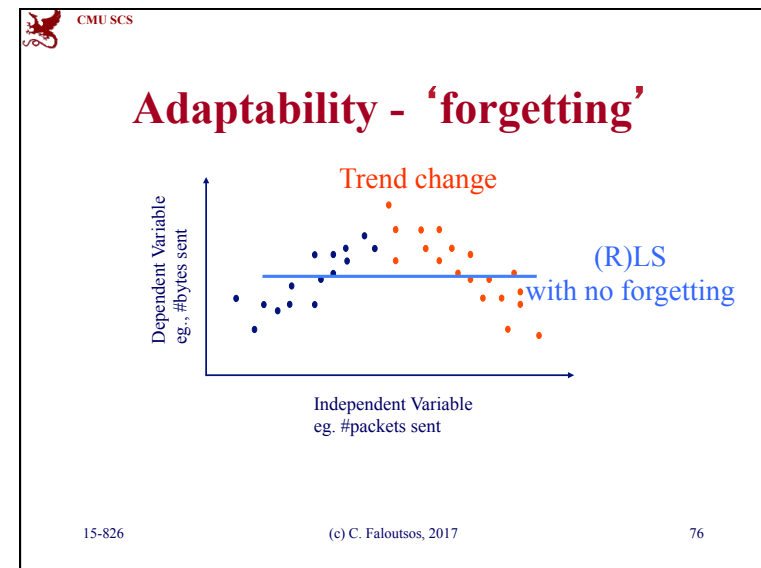




CMU SCS

Even more details

- Q4: can we ‘forget’ the older samples?
- A4: Yes - RLS can easily handle that $[Y_i + \infty]$:

15-826 (c) C. Faloutsos, 2017 74






CMU SCS

Outline

- Motivation
- ...
- Linear Forecasting
 - Auto-regression: Least Squares; RLS
 - Co-evolving time sequences
 - Examples
 - Conclusions

15-826 (c) C. Faloutsos, 2017 89



CMU SCS

Solution:

Q: what should we do?

15-826 (c) C. Faloutsos, 2017 91

CMU SCS

Solution:

Least Squares, with

- Dep. Variable: Repeated(t)
- Indep. Variables: Sent(t-1) ... Sent(t-w);
Lost(t-1) ... Lost(t-w); Repeated(t-1), ...
- (named: 'MUSCLES' [Yi+00])

15-826 (c) C. Faloutsos, 2017 92

CMU SCS

Forecasting - Outline

- Auto-regression
- Least Squares; recursive least squares
- Co-evolving time sequences
- ➔ • Examples
- Conclusions

15-826 (c) C. Faloutsos, 2017 93

CMU SCS

Examples - Experiments

- Datasets
 - Modem pool traffic (14 modems, 1500 time-ticks; #packets per time unit)
 - AT&T WorldNet internet usage (several data streams; 980 time-ticks)
- Measures of success
 - Accuracy : Root Mean Square Error (RMSE)

15-826 (c) C. Faloutsos, 2017 94


CMU SCS

Accuracy - "Modem"

Modem	AR	yesterday	MUSCLES
1	1.8	1.8	1.2
2	1.2	1.2	0.8
3	1.8	1.8	1.2
4	1.8	2.8	1.2
5	1.8	1.8	1.2
6	1.8	2.5	1.2
7	1.8	3.0	1.2
8	1.8	1.8	1.2
9	1.8	2.2	1.2
10	1.8	2.2	1.2
11	1.8	1.8	1.2
12	1.8	2.5	1.2
13	1.8	1.8	1.2
14	1.8	3.8	1.2

MUSCLES outperforms AR & "yesterday"

15-826 (c) C. Faloutsos, 2017 95



CMU SCS

Linear forecasting - Outline

- Auto-regression
- Least Squares; recursive least squares
- Co-evolving time sequences
- Examples
- ➔ Conclusions

15-826 (c) C. Faloutsos, 2017 97

CMU SCS

Conclusions - Practitioner's guide

- AR(IMA) methodology: prevailing method for linear forecasting
- Brilliant method of Recursive Least Squares for fast, incremental estimation.
- See [Box-Jenkins]
- (AWSOM: no human intervention)


15-826 (c) C. Faloutsos, 2017 98

CMU SCS

Resources: software and urls

- free-ware: 'R' for stat. analysis (clone of Splus)
<http://cran.r-project.org/>
- python script for RLS
<http://www.cs.cmu.edu/~christos/SRC/rls-all.tar>

15-826 (c) C. Faloutsos, 2017 99




CMU SCS

Books

- George E.P. Box and Gwilym M. Jenkins and Gregory C. Reinsel, *Time Series Analysis: Forecasting and Control*, Prentice Hall, 1994 (the classic book on ARIMA, 3rd ed.)
- Brockwell, P. J. and R. A. Davis (1987). *Time Series: Theory and Methods*. New York, Springer Verlag.

15-826 (c) C. Faloutsos, 2017 100




CMU SCS

Additional Reading

- [Papadimitriou+ vldb2003] Spiros Papadimitriou, Anthony Brockwell and Christos Faloutsos *Adaptive, Hands-Off Stream Mining* VLDB 2003, Berlin, Germany, Sept. 2003
- [Yi+00] Byoung-Kee Yi et al.: *Online Data Mining for Co-Evolving Time Sequences*, ICDE 2000. (Describes MUSCLES and Recursive Least Squares)

15-826 (c) C. Faloutsos, 2017 101




CMU SCS

Outline

- Motivation
- Similarity search and distance functions
- Linear Forecasting
- ➔ • Bursty traffic - fractals and multifractals
- Non-linear forecasting
- Conclusions

15-826 (c) C. Faloutsos, 2017 102



CMU SCS

Bursty Traffic & Multifractals

15-826 (c) C. Faloutsos, 2017 103

CMU SCS

Outline

- Motivation
- ...
- Linear Forecasting
- Bursty traffic - fractals and multifractals
 - ➔ – Problem
 - Main idea (80/20, Hurst exponent)
 - Results

15-826 (c) C. Faloutsos, 2017 104

CMU SCS

Reference:

[Wang+02] Mengzhi Wang, Tara Madhyastha, Ngai Hang Chang, Spiros Papadimitriou and Christos Faloutsos, *Data Mining Meets Performance Evaluation: Fast Algorithms for Modeling Bursty Traffic*, ICDE 2002, San Jose, CA, 2/26/2002 - 3/1/2002.

Full thesis: CMU-CS-05-185
Performance Modeling of Storage Devices using Machine Learning Mengzhi Wang, Ph.D. Thesis
[Abstract](#), [.ps.gz](#), [.pdf](#)

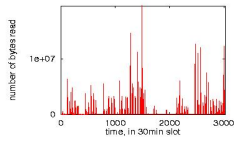
15-826 (c) C. Faloutsos, 2017 105

CMU SCS

Recall: Problem #1:

Goal: given a signal (eg., #bytes over time)
 Find: patterns, periodicities, and/or compress

#bytes



Bytes per 30' (packets per day; earthquakes per year)

time

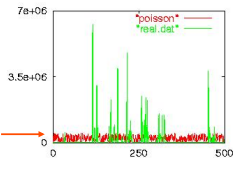
15-826 (c) C. Faloutsos, 2017 106

CMU SCS

Problem #1

- model bursty traffic
- generate realistic traces
- (Poisson does not work)


bytes



Poisson

time

15-826 (c) C. Faloutsos, 2017 107




CMU SCS

Motivation

- predict queue length distributions (e.g., to give probabilistic guarantees)
- “learn” traffic, for buffering, prefetching, ‘active disks’, web servers

15-826 (c) C. Faloutsos, 2017 108




CMU SCS

But:

- Q1: How to generate realistic traces; extrapolate; give guarantees?
- Q2: How to estimate the model parameters?

15-826 (c) C. Faloutsos, 2017 109




CMU SCS

Outline

- Motivation
- ...
- Linear Forecasting
- Bursty traffic - fractals and multifractals
 - Problem
 - ➔ – Main idea (80/20, Hurst exponent)
 - Results

15-826 (c) C. Faloutsos, 2017 110



CMU SCS

Approach

- Q1: How to generate a sequence, that is
 - bursty
 - self-similar
 - and has similar queue length distributions

15-826 (c) C. Faloutsos, 2017 111

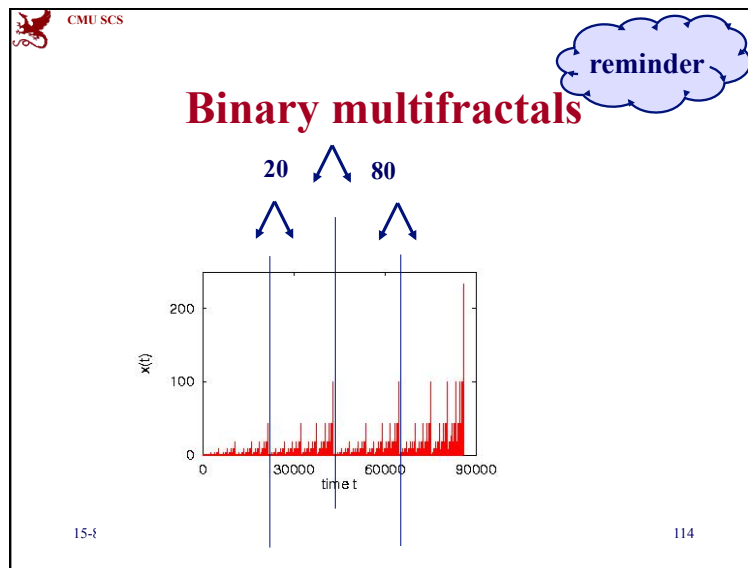

CMU SCS

reminder

Approach

- A: 'binomial multifractal' [Wang+02]
- ~ 80 -20 'law':
 - 80% of bytes/queries etc on first half
 - repeat recursively
- b : bias factor (eg., 80%)

15-826 (c) C. Faloutsos, 2017 112




CMU SCS

Could you use IFS?

To generate such traffic?


15-826 (c) C. Faloutsos, 2017 115



Could you use IFS?

To generate such traffic?
 A: Yes – which transformations?

15-826 (c) C. Faloutsos, 2017 116




Could you use IFS?

To generate such traffic?
 A: Yes – which transformations?
 A:

$$x' = x / 2 \quad (p = 0.2)$$

$$x' = x / 2 + 0.5 \quad (p = 0.8)$$


15-826 (c) C. Faloutsos, 2017 117



Parameter estimation

- Q2: How to estimate the bias factor b ?

15-826 (c) C. Faloutsos, 2017 118



Parameter estimation


- Q2: How to estimate the bias factor b ?
- A: MANY ways [Crovella+96]
 - Hurst exponent
 - variance plot
 - even DFT amplitude spectrum! ('periodogram')
 - Fractal dimension (D2)
 - Or D1 ('entropy plot' [Wang+02])


15-826 (c) C. Faloutsos, 2017 119


CMU SCS

Fractal dimension

- Real (and 80-20) datasets can be in-between: bursts, gaps, smaller bursts, smaller gaps, at every scale

Dim = 1 

Dim=0 

0<Dim<1 

15-826 (c) C. Faloutsos, 2017 133

CMU SCS

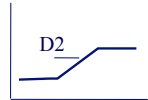
Estimating 'b'

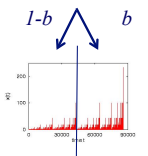
- Exercise:** Show that

$$D_2 = -\log_2 (b^2 + (1-b)^2)$$

Sanity checks:

- $b = 1.0$ $D_2 = ??$
- $b = 0.5$ $D_2 = ??$

Log (#pairs(<r)) 



15-826 (c) C. Faloutsos, 2017 134

CMU SCS

(Fractals, again)

- What set of points could have behavior between point and line?


15-826 (c) C. Faloutsos, 2017 135

CMU SCS

Cantor dust


- Eliminate the middle third
- Recursively!

15-826 (c) C. Faloutsos, 2017 136




CMU SCS

Cantor dust




15-826 (c) C. Faloutsos, 2017 137




CMU SCS

Cantor dust




15-826 (c) C. Faloutsos, 2017 138




CMU SCS

Cantor dust




15-826 (c) C. Faloutsos, 2017 139




CMU SCS


Cantor dust



15-826 (c) C. Faloutsos, 2017 140




Cantor dust



Dimensionality?
(no length; infinite # points!)
Answer: $\log 2 / \log 3 = 0.6$


15-826 (c) C. Faloutsos, 2017 141



Conclusions

- Multifractals (80/20, 'b-model', Multiplicative Wavelet Model (MWM)) for analysis and synthesis of bursty traffic


15-826 (c) C. Faloutsos, 2017 149



Books

- Fractals: Manfred Schroeder: *Fractals, Chaos, Power Laws: Minutes from an Infinite Paradise* W.H. Freeman and Company, 1991 (Probably the BEST book on fractals!)


15-826 (c) C. Faloutsos, 2017 150



Further reading:

- Crovella, M. and A. Bestavros (1996). Self-Similarity in World Wide Web Traffic, Evidence and Possible Causes. *Sigmetrics*.
- [ieeeTN94] W. E. Leland, M.S. Taqqu, W. Willinger, D.V. Wilson, *On the Self-Similar Nature of Ethernet Traffic*, IEEE Transactions on Networking, 2, 1, pp 1-15, Feb. 1994.

15-826 (c) C. Faloutsos, 2017 151


 CMU SCS

Further reading

- [Riedi+99] R. H. Riedi, M. S. Crouse, V. J. Ribeiro, and R. G. Baraniuk, *A Multifractal Wavelet Model with Application to Network Traffic*, IEEE Special Issue on Information Theory, 45. (April 1999), 992-1018.
- [Wang+02] Mengzhi Wang, Tara Madhyastha, Ngai Hang Chang, Spiros Papadimitriou and Christos Faloutsos, *Data Mining Meets Performance Evaluation: Fast Algorithms for Modeling Bursty Traffic*, ICDE 2002, San Jose, CA, 2/26/2002 - 3/1/2002.

Entropy plots


15-826 (c) C. Faloutsos, 2017 152

 CMU SCS

Outline


- Motivation
- ...
- Linear Forecasting
- Bursty traffic - fractals and multifractals
- ➔ • Non-linear forecasting
- Conclusions

15-826 (c) C. Faloutsos, 2017 153

 CMU SCS

Chaos and non-linear forecasting

15-826 (c) C. Faloutsos, 2017 154

 CMU SCS

Reference:

[Deepay Chakrabarti and Christos Faloutsos
*F4: Large-Scale Automated Forecasting
using Fractals* CIKM 2002, Washington
DC, Nov. 2002.]

15-826 (c) C. Faloutsos, 2017 155

CMU SCS

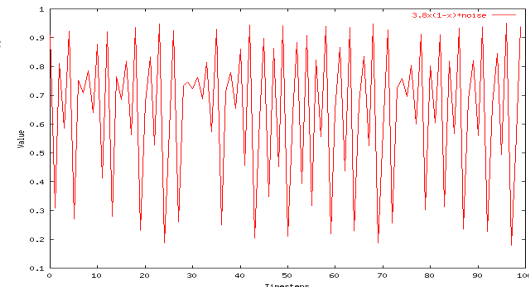
Detailed Outline

- Non-linear forecasting
 - Problem
 - Idea
 - How-to
 - Experiments
 - Conclusions

15-826 (c) C. Faloutsos, 2017 156

CMU SCS

Recall: Problem #1



Value

Time

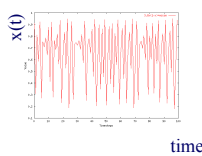

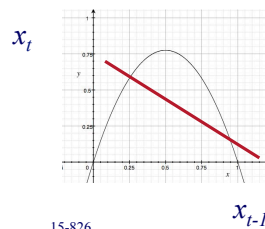
Given a time series $\{x_t\}$, predict its future course, that is, x_{t+1} , x_{t+2} , ...

15-826 (c) C. Faloutsos, 2017 157

CMU SCS

Datasets

Logistic Parabola:
 $x_t = ax_{t-1}(1-x_{t-1}) + \text{noise}$
 Models population of flies [R. May/1976]

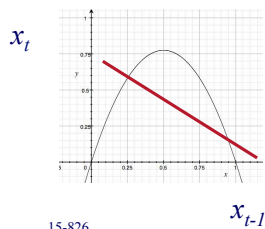
Lag-plot
ARIMA: fails

15-826 (c) C. Faloutsos, 2017 158

CMU SCS

How to forecast?

- ARIMA - but: linearity assumption



Lag-plot
ARIMA: fails


15-826 (c) C. Faloutsos, 2017 159

CMU SCS

How to forecast?

- ARIMA - but: linearity assumption
- ANSWER: ‘Delayed Coordinate Embedding’ = Lag Plots [Sauer92]
~ nearest-neighbor search, for past incidents

15-826 (c) C. Faloutsos, 2017 160



CMU SCS

Questions:

- Q1: How to choose lag L ?
- Q2: How to choose k (the # of NN)?
- Q3: How to interpolate?
- Q4: why should this work at all?

15-826 (c) C. Faloutsos, 2017 162

CMU SCS

Q1: Choosing lag L

- Manually (16, in award winning system by [Sauer94])

15-826 (c) C. Faloutsos, 2017 163

CMU SCS

Q2: Choosing number of neighbors k

- Manually (typically $\sim 1-10$)

15-826 (c) C. Faloutsos, 2017 164

CMU SCS

Q3: How to interpolate?

How do we interpolate between the k nearest neighbors?

A3.1: Average

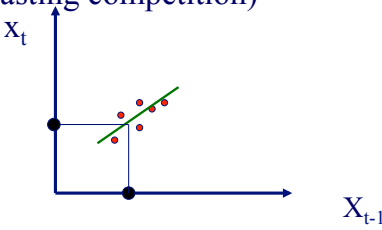
A3.2: Weighted average (weights drop with distance - how?)

15-826 (c) C. Faloutsos, 2017 165

CMU SCS

Q3: How to interpolate?

A3.3: Using SVD - seems to perform best ([Sauer94] - first place in the Santa Fe forecasting competition)



15-826 (c) C. Faloutsos, 2017 166

CMU SCS

Q4: Any theory behind it?

A4: YES!

15-826 (c) C. Faloutsos, 2017 167

CMU SCS

Theoretical foundation

- Based on the ‘Takens theorem’ [Takens81]
- which says that long enough delay vectors **can do prediction**, even if there are unobserved variables in the dynamical system (= diff. equations)

15-826 (c) C. Faloutsos, 2017 168

CMU SCS

Theoretical foundation

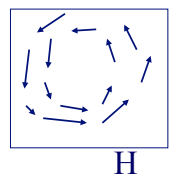
Example: Lotka-Volterra equations

$$\frac{dH}{dt} = rH - aH*P$$

$$\frac{dP}{dt} = bH*P - mP$$

H is count of prey (e.g., hare)
P is count of predators (e.g., lynx)

Suppose only P(t) is observed (t=1, 2, ...).



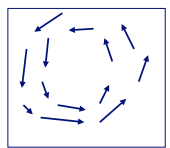
15-826 (c) C. Faloutsos, 2017 169

CMU SCS

Theoretical foundation

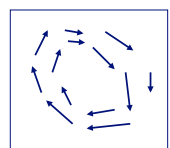
- But the delay vector space is a faithful reconstruction of the internal system state
- So prediction in **delay vector space** is as good as prediction in **state space**

P



H

P(t)



P(t-1)

15-826 (c) C. Faloutsos, 2017 170

CMU SCS

Detailed Outline

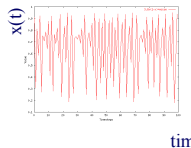

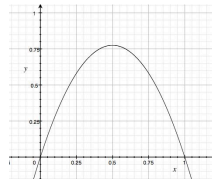
- Non-linear forecasting
 - Problem
 - Idea
 - How-to
 - ➔ – Experiments
 - Conclusions

15-826 (c) C. Faloutsos, 2017 171

CMU SCS

Datasets

Logistic Parabola:
 $x_t = ax_{t-1}(1-x_{t-1}) + \text{noise}$
 Models population of flies [R. May/1976]

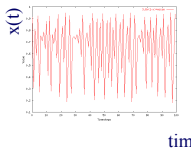
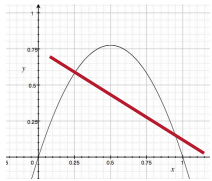
Lag-plot

15-826 (c) C. Faloutsos, 2017 172

CMU SCS

Datasets

Logistic Parabola:
 $x_t = ax_{t-1}(1-x_{t-1}) + \text{noise}$
 Models population of flies [R. May/1976]

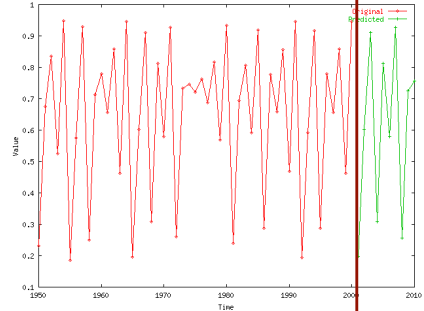
Lag-plot
 ARIMA: fails

15-826 (c) C. Faloutsos, 2017 173

CMU SCS

Logistic Parabola

Value



Time

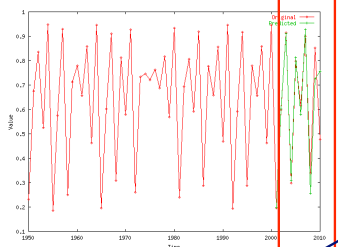
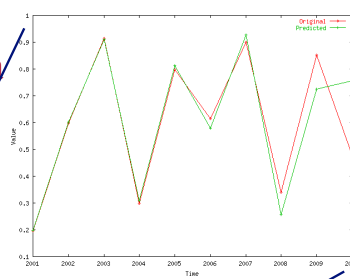
Timesteps

15-826 (c) C. Faloutsos, 2017 174

CMU SCS

Logistic Parabola

Comparison of prediction to correct values


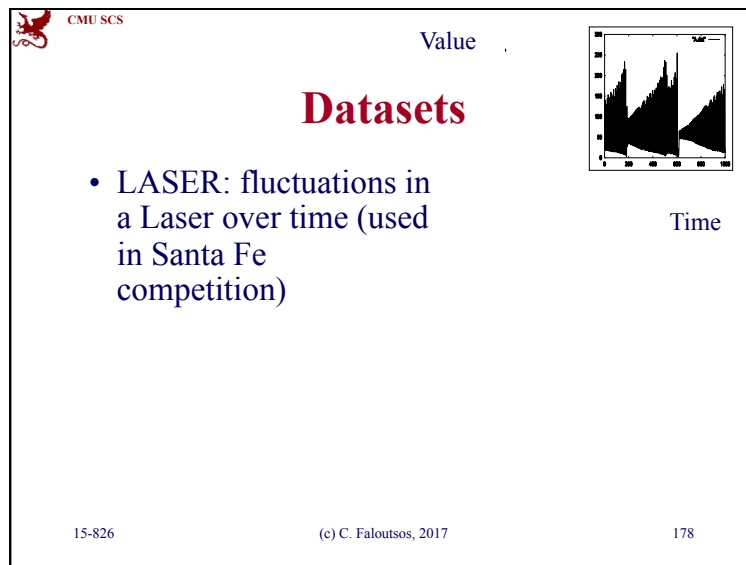






Value

Time

Timesteps

15-826 (c) C. Faloutsos, 2017 175




 CMU SCS

Conclusions

- Lag plots for non-linear forecasting (Takens' theorem)
- suitable for 'chaotic' signals


15-826 (c) C. Faloutsos, 2017 180

 CMU SCS

References

- Deepay Chakrabarti and Christos Faloutsos *F4: Large-Scale Automated Forecasting using Fractals* CIKM 2002, Washington DC, Nov. 2002.
- Sauer, T. (1994). *Time series prediction using delay coordinate embedding*. (in book by Weigend and Gershenfeld, below) Addison-Wesley.
- Takens, F. (1981). *Detecting strange attractors in fluid turbulence*. Dynamical Systems and Turbulence. Berlin: Springer-Verlag.


15-826 (c) C. Faloutsos, 2017 181

 CMU SCS

References

- Weigend, A. S. and N. A. Gerschenfeld (1994). *Time Series Prediction: Forecasting the Future and Understanding the Past*, Addison Wesley. (Excellent collection of papers on chaotic/non-linear forecasting, describing the algorithms behind the winners of the Santa Fe competition.)


15-826 (c) C. Faloutsos, 2017 182

 CMU SCS

Overall conclusions

- Similarity search: **Euclidean**/time-warping; **feature extraction** and **SAMs**

15-826 (c) C. Faloutsos, 2017 183




CMU SCS

Overall conclusions

- Similarity search: **Euclidean**/time-warping; **feature extraction** and **SAMs**
- Signal processing: **DWT** is a powerful tool

15-826 (c) C. Faloutsos, 2017 184




CMU SCS

Overall conclusions

- Similarity search: **Euclidean**/time-warping; **feature extraction** and **SAMs**
- Signal processing: **DWT** is a powerful tool
- Linear Forecasting: **AR** (Box-Jenkins) methodology; **AWSOM**

15-826 (c) C. Faloutsos, 2017 185




CMU SCS

Overall conclusions

- Similarity search: **Euclidean**/time-warping; **feature extraction** and **SAMs**
- Signal processing: **DWT** is a powerful tool
- Linear Forecasting: **AR** (Box-Jenkins) methodology; **AWSOM**
- Bursty traffic: **multifractals** (80-20 'law')

15-826 (c) C. Faloutsos, 2017 186



CMU SCS

Overall conclusions

- Similarity search: **Euclidean**/time-warping; **feature extraction** and **SAMs**
- Signal processing: **DWT** is a powerful tool
- Linear Forecasting: **AR** (Box-Jenkins) methodology; **AWSOM**
- Bursty traffic: **multifractals** (80-20 'law')
- Non-linear forecasting: **lag-plots** (Takens)

15-826 (c) C. Faloutsos, 2017 187