

```
2017-11-20 02:39:53,949 INFO org.apache.hadoop.yarn.client.api.impl.TimelineClientImpl (main):
Timeline service address: http://ip-172-31-20-220.us-east-2.compute.internal:8188/ws/v1/timeline/
2017-11-20 02:39:53,963 INFO org.apache.hadoop.yarn.client.RMPProxy (main): Connecting to
ResourceManager at ip-172-31-20-220.us-east-2.compute.internal/172.31.20.220:8032
2017-11-20 02:39:54,484 INFO org.apache.hadoop.yarn.client.api.impl.TimelineClientImpl (main):
Timeline service address: http://ip-172-31-20-220.us-east-2.compute.internal:8188/ws/v1/timeline/
2017-11-20 02:39:54,485 INFO org.apache.hadoop.yarn.client.RMPProxy (main): Connecting to
ResourceManager at ip-172-31-20-220.us-east-2.compute.internal/172.31.20.220:8032
2017-11-20 02:39:56,521 WARN com.amazon.ws.emr.hadoop.fs.rolemapping.RoleMappings (main): Found no
mappings configured with 'fs.s3.authorization.roleMapping', credentials resolution may not work as
expected
2017-11-20 02:39:58,124 WARN org.apache.hadoop.mapreduce.JobResourceUploader (main): Hadoop command-
line option parsing not performed. Implement the Tool interface and execute your application with
ToolRunner to remedy this.
2017-11-20 02:40:01,039 INFO com.hadoop.compression.lzo.GPLNativeCodeLoader (main): Loaded native gpl
library
2017-11-20 02:40:01,043 INFO com.hadoop.compression.lzo.LzoCodec (main): Successfully loaded &
initialized native-lzo library [hadoop-lzo rev 5f779390180acb8f1d86999c0a0294917976289f]
2017-11-20 02:40:01,121 INFO org.apache.hadoop.mapred.FileInputFormat (main): Total input paths to
process : 3
2017-11-20 02:40:01,200 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): number of splits:9
2017-11-20 02:40:01,381 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): Submitting tokens for
job: job_1511143224443_0002
2017-11-20 02:40:01,811 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl (main): Submitted
application application_1511143224443_0002
2017-11-20 02:40:02,004 INFO org.apache.hadoop.mapreduce.Job (main): The url to track the job:
http://ip-172-31-20-220.us-east-2.compute.internal:20888/proxy/application_1511143224443_0002/
2017-11-20 02:40:02,019 INFO org.apache.hadoop.mapreduce.Job (main): Running job:
job_1511143224443_0002
2017-11-20 02:40:09,216 INFO org.apache.hadoop.mapreduce.Job (main): Job job_1511143224443_0002
running in uber mode : false
2017-11-20 02:40:09,217 INFO org.apache.hadoop.mapreduce.Job (main): map 0% reduce 0%
2017-11-20 02:40:24,397 INFO org.apache.hadoop.mapreduce.Job (main): map 11% reduce 0%
2017-11-20 02:40:25,402 INFO org.apache.hadoop.mapreduce.Job (main): map 22% reduce 0%
2017-11-20 02:40:33,460 INFO org.apache.hadoop.mapreduce.Job (main): map 33% reduce 0%
2017-11-20 02:40:34,465 INFO org.apache.hadoop.mapreduce.Job (main): map 56% reduce 0%
2017-11-20 02:40:36,475 INFO org.apache.hadoop.mapreduce.Job (main): map 78% reduce 0%
2017-11-20 02:40:37,481 INFO org.apache.hadoop.mapreduce.Job (main): map 89% reduce 0%
2017-11-20 02:40:45,516 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 15%
2017-11-20 02:40:46,523 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2017-11-20 02:40:47,535 INFO org.apache.hadoop.mapreduce.Job (main): Job job_1511143224443_0002
completed successfully
2017-11-20 02:40:47,648 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 55
```

#### File System Counters

```
FILE: Number of bytes read=5651
FILE: Number of bytes written=1423967
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=807
HDFS: Number of bytes written=0
HDFS: Number of read operations=9
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
S3: Number of bytes read=43272
S3: Number of bytes written=6604
S3: Number of read operations=0
S3: Number of large read operations=0
S3: Number of write operations=0
```

#### Job Counters

```
Killed map tasks=1
Launched map tasks=9
Launched reduce tasks=2
Data-local map tasks=9
```

Total time spent by all maps in occupied slots (ms)=6485136  
Total time spent by all reduces in occupied slots (ms)=1841952  
Total time spent by all map tasks (ms)=135107  
Total time spent by all reduce tasks (ms)=19187  
Total vcore-milliseconds taken by all map tasks=135107  
Total vcore-milliseconds taken by all reduce tasks=19187  
Total megabyte-milliseconds taken by all map tasks=207524352  
Total megabyte-milliseconds taken by all reduce tasks=58942464

Map-Reduce Framework

Map input records=12  
Map output records=1914  
Map output bytes=37449  
Map output materialized bytes=16755  
Input split bytes=807  
Combine input records=0  
Combine output records=0  
Reduce input groups=349  
Reduce shuffle bytes=16755  
Reduce input records=1914  
Reduce output records=349  
Spilled Records=3828  
Shuffled Maps =18  
Failed Shuffles=0  
Merged Map outputs=18  
GC time elapsed (ms)=4884  
CPU time spent (ms)=10810  
Physical memory (bytes) snapshot=5586518016  
Virtual memory (bytes) snapshot=39040970752  
Total committed heap usage (bytes)=5028970496

Shuffle Errors

BAD\_ID=0  
CONNECTION=0  
IO\_ERROR=0  
WRONG\_LENGTH=0  
WRONG\_MAP=0  
WRONG\_REDUCE=0

File Input Format Counters

Bytes Read=43272

File Output Format Counters

Bytes Written=6604