
 CMU SCS

15-826: Multimedia Databases and Data Mining

Lecture#5: Multi-key and Spatial Access Methods – II – z-ordering
C. Faloutsos

 CMU SCS

Must-read material

- MM-Textbook, Chapter 5.1
- Ramakrishnan+Gehrke, Chapter 28.4
- J. Orenstein, [*Spatial Query Processing in an Object-Oriented Database System*](#), Proc. ACM SIGMOD, May, 1986, pp. 326-336, Washington D.C.

15-826 Copyright: C. Faloutsos (2017) 2


 CMU SCS

Outline

Goal: ‘Find similar / interesting things’

- Intro to DB
- ➔ • Indexing - similarity search
- Data Mining

15-826 Copyright: C. Faloutsos (2017) 3

 CMU SCS

Indexing - Detailed outline

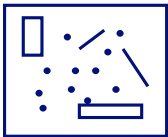
- primary key indexing
- secondary key / multi-key indexing
- ➔ • spatial access methods
 - problem defn
 - z-ordering
 - R-trees
 - ...
- text
- ...

15-826 Copyright: C. Faloutsos (2017) 4

CMU SCS

Spatial Access Methods - problem

- Given a collection of geometric objects (points, lines, polygons, ...)
- organize them on disk, to answer spatial queries (like??)

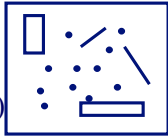


15-826 Copyright: C. Faloutsos (2017) 5

CMU SCS

Spatial Access Methods - problem

- Given a collection of geometric objects (points, lines, polygons, ...)
- organize them on disk, to answer
 - point queries
 - range queries
 - k-nn queries
 - spatial joins (‘all pairs’ queries)

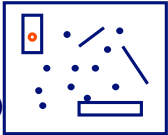


15-826 Copyright: C. Faloutsos (2017) 6

CMU SCS

Spatial Access Methods - problem

- Given a collection of geometric objects (points, lines, polygons, ...)
- organize them on disk, to answer
 - point queries
 - range queries
 - k-nn queries
 - spatial joins (‘all pairs’ queries)

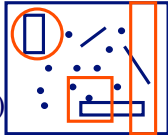


15-826 Copyright: C. Faloutsos (2017) 7

CMU SCS

Spatial Access Methods - problem

- Given a collection of geometric objects (points, lines, polygons, ...)
- organize them on disk, to answer
 - point queries
 - range queries
 - k-nn queries
 - spatial joins (‘all pairs’ queries)

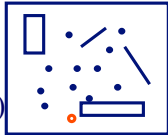


15-826 Copyright: C. Faloutsos (2017) 8

CMU SCS

Spatial Access Methods - problem

- Given a collection of geometric objects (points, lines, polygons, ...)
- organize them on disk, to answer
 - point queries
 - range queries
 - k-nn queries
 - spatial joins ('all pairs' queries)

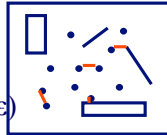


15-826 Copyright: C. Faloutsos (2017) 9

CMU SCS

Spatial Access Methods - problem

- Given a collection of geometric objects (points, lines, polygons, ...)
- organize them on disk, to answer
 - point queries
 - range queries
 - k-nn queries
 - spatial joins ('all pairs' within ϵ)



15-826 Copyright: C. Faloutsos (2017) 10

CMU SCS

SAMs - motivation

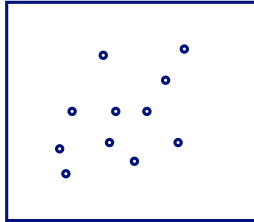
- Q: applications?

15-826 Copyright: C. Faloutsos (2017) 11

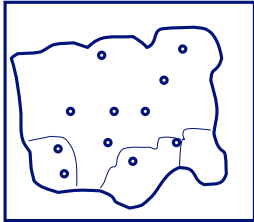
CMU SCS

SAMs - motivation

traditional DB

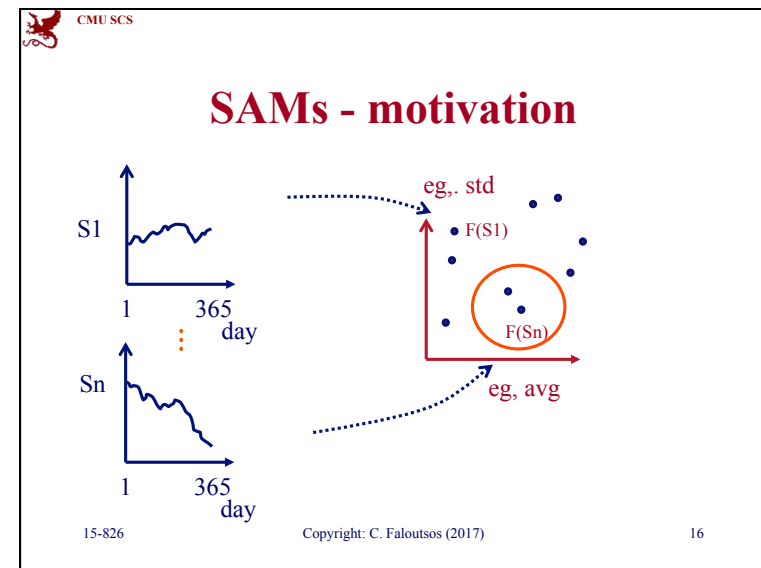
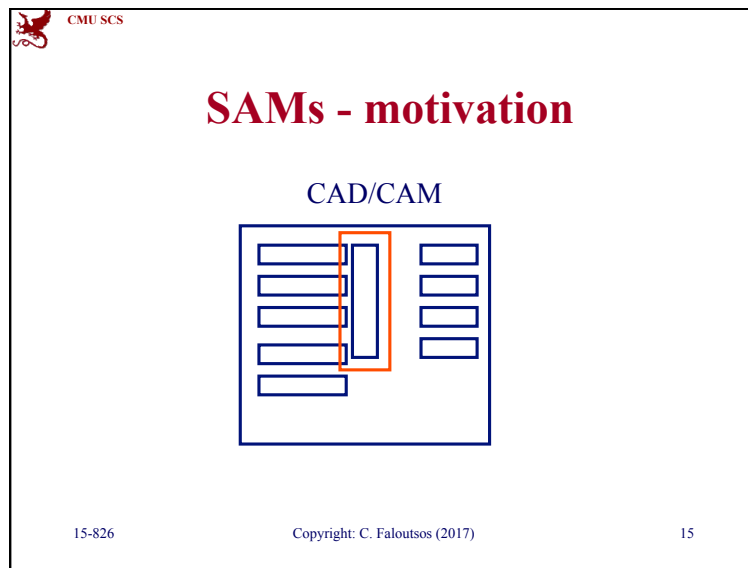
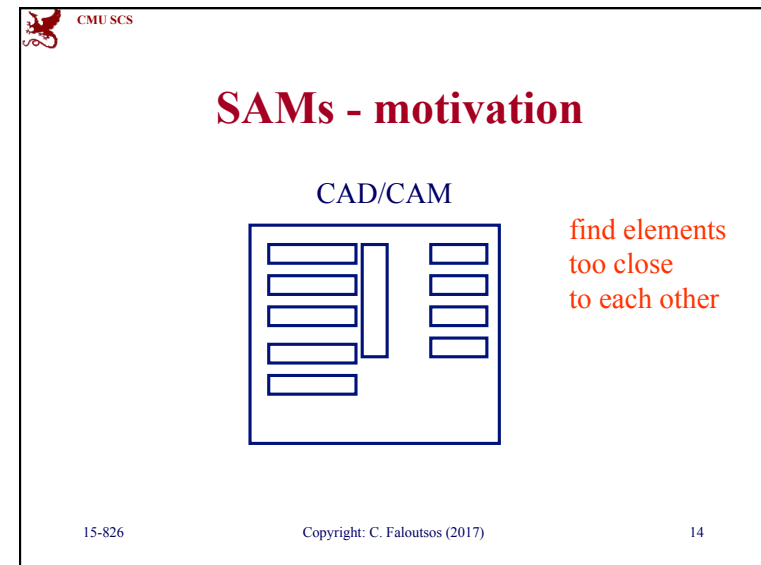
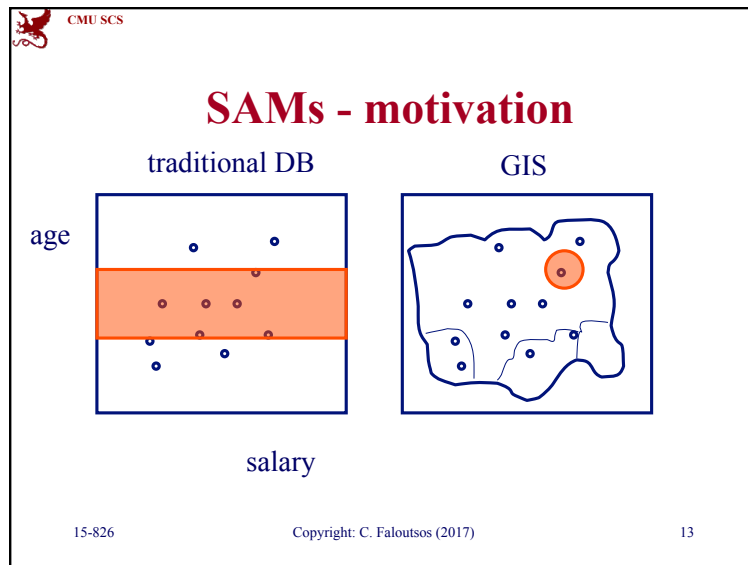


GIS



salary

15-826 Copyright: C. Faloutsos (2017) 12



CMU SCS

Indexing - Detailed outline

- primary key indexing
- secondary key / multi-key indexing
- spatial access methods
 - problem defn
 - ➔ – z-ordering
 - R-trees
 - ...
- text
- ...

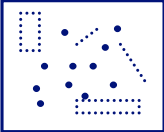
15-826 Copyright: C. Faloutsos (2017) 17

CMU SCS

SAMs: solutions

- z-ordering
- R-trees
- (grid files)

Q: how would you organize, e.g., n -dim points, on disk? (C points per disk page)



15-826 Copyright: C. Faloutsos (2017) 18

CMU SCS

z-ordering

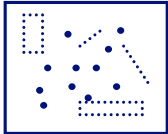
Q: how would you organize, e.g., n -dim points, on disk? (C points per disk page)

Hint: reduce the problem to 1-d points (!!)

Q1: why?

A:

Q2: how?



15-826 Copyright: C. Faloutsos (2017) 19

CMU SCS

z-ordering

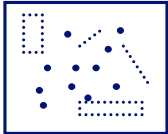
Q: how would you organize, e.g., n -dim points, on disk? (C points per disk page)

Hint: reduce the problem to 1-d points (!!)

Q1: why?

A: B-trees!

Q2: how?




15-826 Copyright: C. Faloutsos (2017) 20

z-ordering

Q2: how?

A: assume finite granularity; z-ordering = bit-shuffling = N-trees = Morton keys = geocoding = ...



15-826

Copyright: C. Faloutsos (2017)

21

z-ordering

Q2: how?

A: assume finite granularity (e.g., $2^{32} \times 2^{32}$; 4x4 here)

Q2.1: how to map n-d cells to 1-d cells?




15-826

Copyright: C. Faloutsos (2017)

22

z-ordering

Q2.1: how to map n -d cells to 1-d cells?



15-826

Copyright: C. Faloutsos (2017)


23

z-ordering

Q2.1: how to map n -d cells to 1-d cells?

A: row-wise

Q: is it good?



15-826

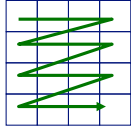
Copyright: C. Faloutsos (2017)

24

CMU SCS

z-ordering

Q: is it good?
A: great for 'x' axis; bad for 'y' axis

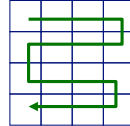


15-826 Copyright: C. Faloutsos (2017) 25

CMU SCS

z-ordering

Q: How about the 'snake' curve?

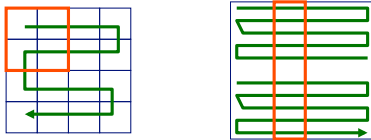


15-826 Copyright: C. Faloutsos (2017) 26

CMU SCS

z-ordering

Q: How about the 'snake' curve?
A: still problems:



2^{32}

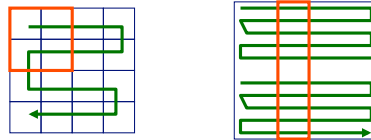
2^{32}

15-826 Copyright: C. Faloutsos (2017) 27

CMU SCS

z-ordering

Q: Why are those curves 'bad' ?
A: no distance preservation (~ clustering)
Q: solution?



2^{32}

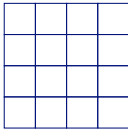
2^{32}

15-826 Copyright: C. Faloutsos (2017) 28

CMU SCS

z-ordering

Q: solution? (w/ good clustering, and easy to compute, for 2-d and n -d?)



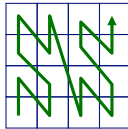
15-826 Copyright: C. Faloutsos (2017) 29

CMU SCS

z-ordering

Q: solution? (w/ good clustering, and easy to compute, for 2-d and n -d?)

A: z-ordering/bit-shuffling/linear-quadtrees



‘looks’ better:

- few long jumps;
- scoops out the whole quadrant before leaving it
- a.k.a. space filling curves

15-826 Copyright: C. Faloutsos (2017) 30

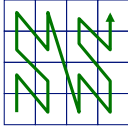
CMU SCS

z-ordering

z-ordering/bit-shuffling/linear-quadtrees

Q: How to generate this curve ($z = f(x,y)$)?

A: 3 (equivalent) answers!



15-826 Copyright: C. Faloutsos (2017) 31

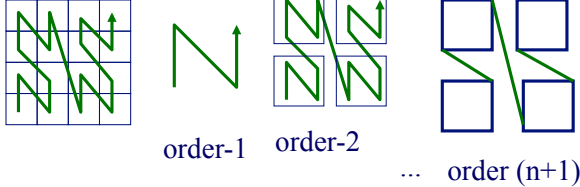
CMU SCS

z-ordering

z-ordering/bit-shuffling/linear-quadtrees

Q: How to generate this curve ($z = f(x,y)$)?

A1: ‘z’ (or ‘N’) shapes, RECURSIVELY



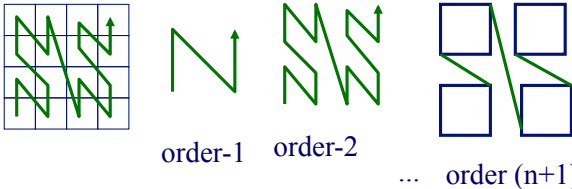
order-1 order-2 ... order (n+1)

15-826 Copyright: C. Faloutsos (2017) 32

z-ordering

Notice:

- self similar (we'll see about fractals, soon)
- method is hard to use: $z = ? f(x,y)$



order-1 order-2 ... order (n+1)

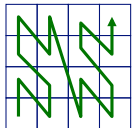
15-826 Copyright: C. Faloutsos (2017) 33

z-ordering

z-ordering/**bit-shuffling**/linear-quadtrees

Q: How to generate this curve ($z = f(x,y)$)?

A: 3 (equivalent) answers!

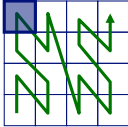


Method #2?

15-826 Copyright: C. Faloutsos (2017) 34

z-ordering

bit-shuffling



x y
 0 0 1 1

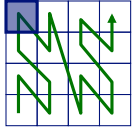
y
 11
 10
 01
 00

00 01 10 11 x

15-826 Copyright: C. Faloutsos (2017) 35

z-ordering

bit-shuffling



x y
 0 0 1 1

y
 11
 10
 01
 00

00 01 10 11 x

$z = (0101)_2 = 5$

15-826 Copyright: C. Faloutsos (2017) 36

CMU SCS

z-ordering

bit-shuffling

How about the reverse:
 $(x,y) = g(z)$?

$$z = (0101)_2 = 5$$

15-826 Copyright: C. Faloutsos (2017) 37

CMU SCS

z-ordering

bit-shuffling

How about n -d spaces?

$$z = (0101)_2 = 5$$

15-826 Copyright: C. Faloutsos (2017) 38

CMU SCS

z-ordering

z-ordering/bit-shuffling/**linear-quadtrees**

Q: How to generate this curve ($z = f(x,y)$)?

A: 3 (equivalent) answers!

Method #3?

15-826 Copyright: C. Faloutsos (2017) 39

CMU SCS

z-ordering

linear-quadtrees : assign N->1, S->0 e.t.c.

15-826 Copyright: C. Faloutsos (2017) 40

z-ordering

... and repeat recursively. Eg.: $z_{\text{blue-cell}} =$
 $WN;WN = (0101)_2 = 5$

15-826 Copyright: C. Faloutsos (2017) 41

z-ordering

Drill: z-value of magenta cell, with the three methods?

15-826 Copyright: C. Faloutsos (2017) 42

z-ordering

Drill: z-value of magenta cell, with the three methods?

method#1: 14
 method#2: $\text{shuffle}(11;10) = (1110)_2 = 14$

15-826 Copyright: C. Faloutsos (2017) 43

z-ordering

Drill: z-value of magenta cell, with the three methods?

method#1: 14
 method#2: $\text{shuffle}(11;10) = (1110)_2 = 14$
 method#3: $EN;ES = \dots = 14$

15-826 Copyright: C. Faloutsos (2017) 44

CMU SCS

z-ordering - Detailed outline

- spatial access methods
 - z-ordering
 - main idea - 3 methods
 - use w/ B-trees; algorithms (range, knn queries ...)
 - non-point (eg., region) data
 - analysis; variations
 - R-trees
 - ...

15-826 Copyright: C. Faloutsos (2017) 45

CMU SCS

z-ordering - usage & algo's

Q1: How to store on disk?
A:
Q2: How to answer range queries etc

15-826 Copyright: C. Faloutsos (2017) 46

CMU SCS

z-ordering - usage & algo's

Q1: How to store on disk?
A: treat z-value as primary key; feed to B-tree

PGH

SF

z	cname	etc
5	SF	
12	PGH	

15-826 Copyright: C. Faloutsos (2017) 47

CMU SCS

z-ordering - usage & algo's

MAJOR ADVANTAGES w/ B-tree:

- already inside commercial systems (no coding/debugging!)
- concurrency & recovery is ready

PGH

SF

z	cname	etc
5	SF	
12	PGH	

15-826 Copyright: C. Faloutsos (2017) 48

CMU SCS

z-ordering - usage & algo's

Q2: queries? (eg.: *find city at (0,3)*)?

PGH

SF

z	cname	etc
5	SF	
12	PGH	

15-826 Copyright: C. Faloutsos (2017) 49

CMU SCS

z-ordering - usage & algo's

Q2: queries? (eg.: *find city at (0,3)*)?

A: find z-value; search B-tree

PGH

SF

z	cname	etc
5	SF	
12	PGH	

15-826 Copyright: C. Faloutsos (2017) 50

CMU SCS

z-ordering - usage & algo's

Q2: range queries?

PGH

SF

z	cname	etc
5	SF	
12	PGH	

15-826 Copyright: C. Faloutsos (2017) 51

CMU SCS

z-ordering - usage & algo's

Q2: range queries?

A: compute ranges of z-values; use B-tree

PGH

SF

9,11-15

z	cname	etc
5	SF	
12	PGH	

15-826 Copyright: C. Faloutsos (2017) 52

CMU SCS

z-ordering - usage & algo's

Q2' : range queries - how to reduce # of qualifying of ranges?

PGH

SF

9,11-15

z	cname	etc
5	SF	
12	PGH	

15-826 Copyright: C. Faloutsos (2017) 53

CMU SCS

z-ordering - usage & algo's

Q2' : range queries - how to reduce # of qualifying of ranges?

A: Augment the query!

PGH

SF

9,11-15 -> 8-15

z	cname	etc
5	SF	
12	PGH	

15-826 Copyright: C. Faloutsos (2017) 54

CMU SCS

z-ordering - usage & algo's

Q2'' : range queries - how to break a query into ranges?

9,11-15

15-826 Copyright: C. Faloutsos (2017) 55

CMU SCS

z-ordering - usage & algo's

Q2'' : range queries - how to break a query into ranges?

A: recursively, quadtree-style; decompose only non-full quadrants

12-15

9,11-15

15-826 Copyright: C. Faloutsos (2017) 56

CMU SCS

z-ordering - usage & algo's

Q2'': range queries - how to break a query into ranges?

A: recursively, quadtree-style; decompose only non-full quadrants

15-826 Copyright: C. Faloutsos (2017) 57

CMU SCS

z-ordering - Detailed outline

- spatial access methods
 - z-ordering
 - main idea - 3 methods
 - use w/ B-trees; algorithms (range, knn queries ...)
 - non-point (eg., region) data
 - analysis; variations
 - R-trees
 - ...

15-826 Copyright: C. Faloutsos (2017) 58

CMU SCS

z-ordering - usage & algo's

Q3: k-nn queries? (say, 1-nn)?

A: traverse B-tree; find nn wrt z-values and ...

z	cname	etc
5	SF	
12	PGH	

15-826 Copyright: C. Faloutsos (2017) 59

CMU SCS

z-ordering - usage & algo's

Q3: k-nn queries? (say, 1-nn)?

A: traverse B-tree; find nn wrt z-values and ...

z	cname	etc
5	SF	
12	PGH	

15-826 Copyright: C. Faloutsos (2017) 60

CMU SCS

z-ordering - usage & algo's

... ask a range query.

15-826 Copyright: C. Faloutsos (2017) 61

CMU SCS

z-ordering - usage & algo's

... ask a range query.

15-826 Copyright: C. Faloutsos (2017) 62

CMU SCS

z-ordering - usage & algo's

Q4: all-pairs queries? (*all pairs of cities within 10 miles from each other?*)

(we'll see 'spatial joins' later: *find all PA counties that intersect a lake*)

15-826 Copyright: C. Faloutsos (2017) 63

CMU SCS

z-ordering - Detailed outline

- spatial access methods
 - z-ordering
 - main idea - 3 methods
 - use w/ B-trees; algorithms (range, knn queries ...)
 - non-point (eg., region) data
 - analysis; variations
 - R-trees
 - ...

15-826 Copyright: C. Faloutsos (2017) 64

z-ordering - regions

Q: z-value for a region?

$z_B = ??$
 $z_C = ??$

15-826 Copyright: C. Faloutsos (2017) 65

z-ordering - regions

Q: z-value for a region?

A: 1 or more z-values; by quadtree decomposition

$z_B = ??$
 $z_C = ??$

15-826 Copyright: C. Faloutsos (2017) 66

z-ordering - regions “don't care”

Q: z-value for a region?

$z_B = 11^{**}$
 $z_C = ??$

15-826 Copyright: C. Faloutsos (2017) 67

z-ordering - regions “don't care”

Q: z-value for a region?

$z_B = 11^{**}$
 $z_C = \{0010; 1000\}$

15-826 Copyright: C. Faloutsos (2017) 68

z-ordering - regions

Q: How to store in B-tree?
Q: How to search (range etc queries)

15-826 Copyright: C. Faloutsos (2017) 69

z-ordering - regions

Q: How to store in B-tree? A: sort ($* < 0 < 1$)
Q: How to search (range etc queries)

z	obj-id	etc
0010	C	
0101	A	
1000	C	
11**	B	

15-826 Copyright: C. Faloutsos (2017) 70

z-ordering - regions

Q: How to search (range etc queries) - eg 'red' range query

z	obj-id	etc
0010	C	
0101	A	
1000	C	
11**	B	

15-826 Copyright: C. Faloutsos (2017) 71

z-ordering - regions

Q: How to search (range etc queries) - eg 'red' range query
A: break query in z-values; check B-tree

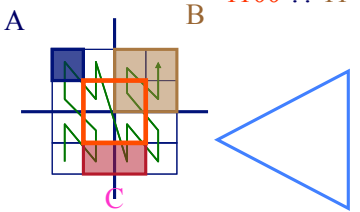
z	obj-id	etc
0010	C	
0101	A	
1000	C	
11**	B	

15-826 Copyright: C. Faloutsos (2017) 72

z-ordering - regions

Almost identical to range queries for point data, except for the “don’t cares” - i.e.,

1100 ?? 11**



z	obj-id	etc
0010	C	
0101	A	
1000	C	
11**	B	

15-826 Copyright: C. Faloutsos (2017) 73

z-ordering - regions

Almost identical to range queries for point data, except for the “don’t cares” - i.e.,

$z1 = 1100 \text{ ?? } 11^{**} = z2$

Specifically: does $z1$ contain/avoid/intersect $z2$?

Q: what is the criterion to decide?

15-826 Copyright: C. Faloutsos (2017) 74

z-ordering - regions

$z1 = 1100 \text{ ?? } 11^{**} = z2$

Specifically: does $z1$ contain/avoid/intersect $z2$?

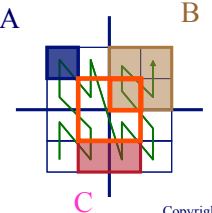
Q: what is the criterion to decide?

A: **Prefix property**: let $r1, r2$ be the corresponding regions, and let $r1$ be the smallest ($\Rightarrow z1$ has fewest ‘*’s). Then:

15-826 Copyright: C. Faloutsos (2017) 75

z-ordering - regions

- $r2$ will either contain completely, or avoid completely $r1$.
- it will contain $r1$, if $z2$ is the prefix of $z1$



1100 ?? 11**

region of $z1$:
completely contained in
region of $z2$

15-826 Copyright: C. Faloutsos (2017) 76

CMU SCS

z-ordering - regions

Drill (True/False). Given:

- $z1 = 011001**$
- $z2 = 01*****$
- $z3 = 0100****$

T/F $r2$ contains $r1$
 T/F $r3$ contains $r1$
 T/F $r3$ contains $r2$

15-826 Copyright: C. Faloutsos (2017) 77

CMU SCS

z-ordering - regions

Drill (True/False). Given:

- $z1 = 011001**$
- $z2 = 01*****$
- $z3 = 0100****$

T/F $r2$ contains $r1$ - TRUE (prefix property)
 T/F $r3$ contains $r1$ - FALSE (disjoint)
 T/F $r3$ contains $r2$ - FALSE ($r2$ contains $r3$)

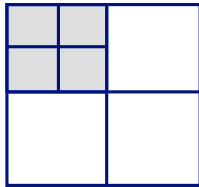
15-826 Copyright: C. Faloutsos (2017) 78

CMU SCS

z-ordering - regions

Drill (True/False). Given:

- $z1 = 011001**$
- $z2 = 01*****$
- $z3 = 0100****$



$z2$

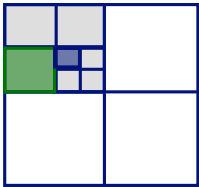
15-826 Copyright: C. Faloutsos (2017) 79

CMU SCS

z-ordering - regions

Drill (True/False). Given:

- $z1 = 011001**$
- $z2 = 01*****$
- $z3 = 0100****$



$z2$

$z3$


T/F $r2$ contains $r1$ - TRUE (prefix property)
 T/F $r3$ contains $r1$ - FALSE (disjoint)
 T/F $r3$ contains $r2$ - FALSE ($r2$ contains $r3$)

15-826 Copyright: C. Faloutsos (2017) 80

CMU SCS

z-ordering - regions

Spatial joins: find (quickly) all
counties intersecting **lakes**

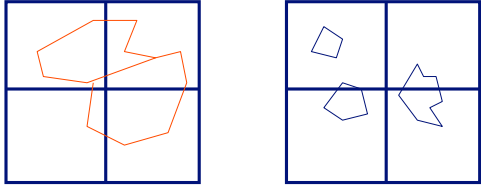


15-826 Copyright: C. Faloutsos (2017) 81

CMU SCS

z-ordering - regions

Spatial joins: find (quickly) all
counties intersecting **lakes**

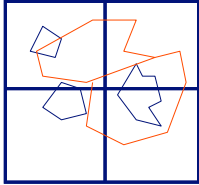


15-826 Copyright: C. Faloutsos (2017) 82

CMU SCS

z-ordering - regions

Spatial joins: find (quickly) all
counties intersecting **lakes**



15-826 Copyright: C. Faloutsos (2017) 83

CMU SCS

z-ordering - regions

Spatial joins: find (quickly) all
counties intersecting **lakes**

Naive algorithm: $O(N * M)$
 Something faster?

15-826 Copyright: C. Faloutsos (2017) 84



CMU SCS

z-ordering - regions

Spatial joins: find (quickly) all
counties intersecting lakes

z	obj-id	etc
0010	ALG	
...	...	
1000	WAS	
11**	ALG	

z	obj-id	etc
0011	Erie	
0101	Erie	
...		
10**	Ont.	

15-826

Copyright: C. Faloutsos (2017)

85



CMU SCS

z-ordering - regions

Spatial joins: find (quickly) all
counties intersecting lakes

Solution: merge the lists of (sorted) z-values,
looking for the prefix property

footnote#1: '*' needs careful treatment
footnote#2: need dup. elimination

15-826

Copyright: C. Faloutsos (2017)

86



CMU SCS

z-ordering - Detailed outline

- spatial access methods
 - z-ordering
 - main idea - 3 methods
 - use w/ B-trees; algorithms (range, knn queries ...)
 - non-point (eg., region) data
 - analysis; variations
 - R-trees
 - ...



15-826

Copyright: C. Faloutsos (2017)


87



CMU SCS

z-ordering - variations

Q: is z-ordering the best we can do?



15-826

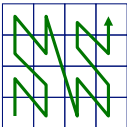
Copyright: C. Faloutsos (2017)

88

CMU SCS

z-ordering - variations

Q: is z-ordering the best we can do?
 A: probably not - occasional long 'jumps'
 Q: then?

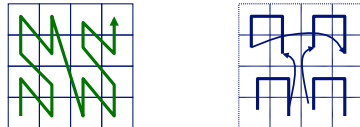


15-826 Copyright: C. Faloutsos (2017) 89

CMU SCS

z-ordering - variations

Q: is z-ordering the best we can do?
 A: probably not - occasional long 'jumps'
 Q: then? A1: Gray codes



15-826 Copyright: C. Faloutsos (2017) 90

CMU SCS

(Gray codes)

- Ingenious way to spot flickering LED – binary:

	000	0
	001	1
	010	2
	011	3
3.5V →	100	4
	101	5
	110	6
	111	7

F. Gray. *Pulse code communication*,
 March 17, 1953
[U.S. Patent 2,632,058](#)

15-826 Copyright: C. Faloutsos (2017) 91

CMU SCS

(Gray codes)

- Ingenious way to spot flickering LED

0
1

15-826 Copyright: C. Faloutsos (2017) 92

CMU SCS

(Gray codes)

- Ingenious way to spot flickering LED

0	.0
1	.1
	..
	..

15-826 Copyright: C. Faloutsos (2017) 93

CMU SCS

(Gray codes)

- Ingenious way to spot flickering LED

0	.0
1	.1
	..
	..

15-826 Copyright: C. Faloutsos (2017) 94

CMU SCS

(Gray codes)

- Ingenious way to spot flickering LED

0	.0
1	.1
	.1
	.0

15-826 Copyright: C. Faloutsos (2017) 95

CMU SCS

(Gray codes)

- Ingenious way to spot flickering LED

0	00
1	01
	11
	10

15-826 Copyright: C. Faloutsos (2017) 96

CMU SCS

(Gray codes)

- Ingenious way to spot flickering LED

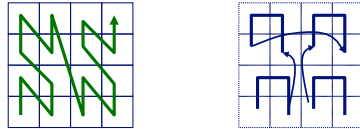
0	00	000	0
1	01	001	1
	11	011	2
	10	010	3
		110	4
		111	5
		101	6
		100	7

15-826 Copyright: C. Faloutsos (2017) 97

CMU SCS

z-ordering - variations

Q: is z-ordering the best we can do?
 A: probably not - occasional long 'jumps'
 Q: then? A1: Gray codes – CAN WE DO BETTER?




15-826 Copyright: C. Faloutsos (2017) 98

CMU SCS

z-ordering - variations

A2: Hilbert curve! (a.k.a. Hilbert-Peano curve)



15-826 Copyright: C. Faloutsos (2017) 99

CMU SCS

(break)



David Hilbert
(1862-1943)




Giuseppe Peano
(1858-1932)

15-826 Copyright: C. Faloutsos (2017) 100

CMU SCS

z-ordering - variations

'Looks' better (never long jumps). How to derive it?

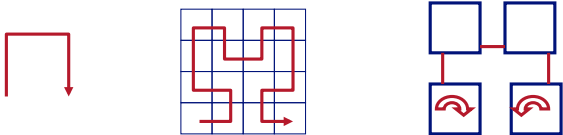


15-826 Copyright: C. Faloutsos (2017) 101

CMU SCS

z-ordering - variations

'Looks' better (never long jumps). How to derive it?



order-1 order-2 ... order (n+1)

15-826 Copyright: C. Faloutsos (2017) 102

CMU SCS

z-ordering - variations

Q: function for the Hilbert curve ($h = f(x,y)$)?

A: bit-shuffling, followed by post-processing, to account for rotations. Linear on # bits. See textbook, for pointers to code/ algorithms (eg., [Jagadish, 90])

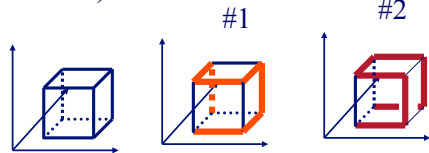
15-826 Copyright: C. Faloutsos (2017) 103

CMU SCS

z-ordering - variations

Q: how about Hilbert curve in 3-d? n-d?

A: Exists (and is not unique!). Eg., 3-d, order-1 Hilbert curves (Hamiltonian paths on cube)



#1 #2

15-826 Copyright: C. Faloutsos (2017) 104

CMU SCS

z-ordering - Detailed outline

- spatial access methods
 - z-ordering
 - main idea - 3 methods
 - use w/ B-trees; algorithms (range, knn queries ...)
 - non-point (eg., region) data
 - analysis; variations
 - R-trees
 - ...

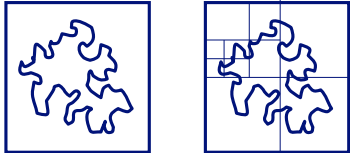
15-826 Copyright: C. Faloutsos (2017) 105

CMU SCS

z-ordering - analysis

Q: How many pieces ('quad-tree blocks') per region?

A: proportional to perimeter (surface etc)

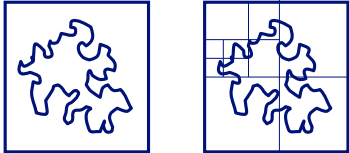


15-826 Copyright: C. Faloutsos (2017) 106

CMU SCS

z-ordering - analysis

(How long is the coastline, say, of England?
Paradox: The answer changes with the yardstick -> fractals ...)

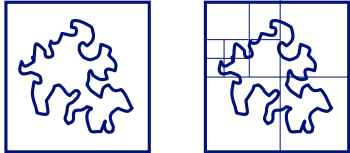


15-826 Copyright: C. Faloutsos (2017) 107

CMU SCS

z-ordering - analysis

Q: Should we decompose a region to full detail (and store in B-tree)?



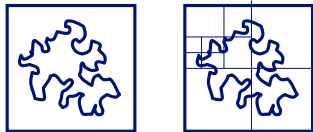
15-826 Copyright: C. Faloutsos (2017) 108

CMU SCS

z-ordering - analysis

Q: Should we decompose a region to full detail (and store in B-tree)?

A: NO! approximation with 1-3 pieces/z-values is best [Orenstein90]




15-826 Copyright: C. Faloutsos (2017) 109

CMU SCS

z-ordering - analysis

Q: how to measure the 'goodness' of a curve?




15-826 Copyright: C. Faloutsos (2017) 110

CMU SCS

z-ordering - analysis

Q: how to measure the 'goodness' of a curve?

A: e.g., avg. # of runs, for range queries



4 runs

(#runs ~ #disk accesses on B-tree)

3 runs

15-826 Copyright: C. Faloutsos (2017) 111

CMU SCS

z-ordering - analysis

Q: So, is Hilbert really better?

A: 27% fewer runs, for 2-d (similar for 3-d)

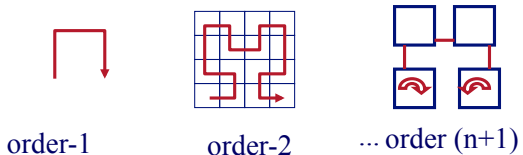
Q: are there formulas for #runs, #of quadtree blocks etc?

A: Yes ([Jagadish; Moon+ etc] see textbook)

15-826 Copyright: C. Faloutsos (2017) 112

z-ordering - fun observations

Hilbert and z-ordering curves: “space filling curves”: eventually, they visit every point in n -d space - therefore:

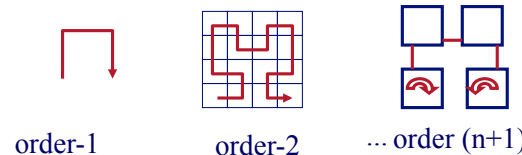


order-1 order-2 ... order (n+1)

15-826 Copyright: C. Faloutsos (2017) 113

z-ordering - fun observations

... they show that the plane has as many points as a line (\rightarrow headaches for 1900's mathematics/topology). (fractals, again!)




order-1 order-2 ... order (n+1)

15-826 Copyright: C. Faloutsos (2017) 114

z-ordering - fun observations

Observation #2: Hilbert (like) curve for video encoding [Y. Matias+, CRYPTO '87]:
Given a frame, visit its pixels in randomized hilbert order; compress; and transmit




15-826 Copyright: C. Faloutsos (2017) 115

z-ordering - fun observations

In general, Hilbert curve is great for preserving distances, clustering, vector quantization etc

15-826 Copyright: C. Faloutsos (2017) 116




CMU SCS

Indexing - Detailed outline

- primary key indexing
- secondary key / multi-key indexing
- spatial access methods
 - problem defn
 - z-ordering
 - ➔ – R-trees
 - ...
- Text

15-826 Copyright: C. Faloutsos (2017) 117



CMU SCS

Conclusions

- z-ordering is a great idea (n-d points -> 1-d points; feed to B-trees)
- used by TIGER system
<http://www.census.gov/geo/www/tiger/>
- and (most probably) by other GIS products
- works great with low-dim points

15-826 Copyright: C. Faloutsos (2017) 118